





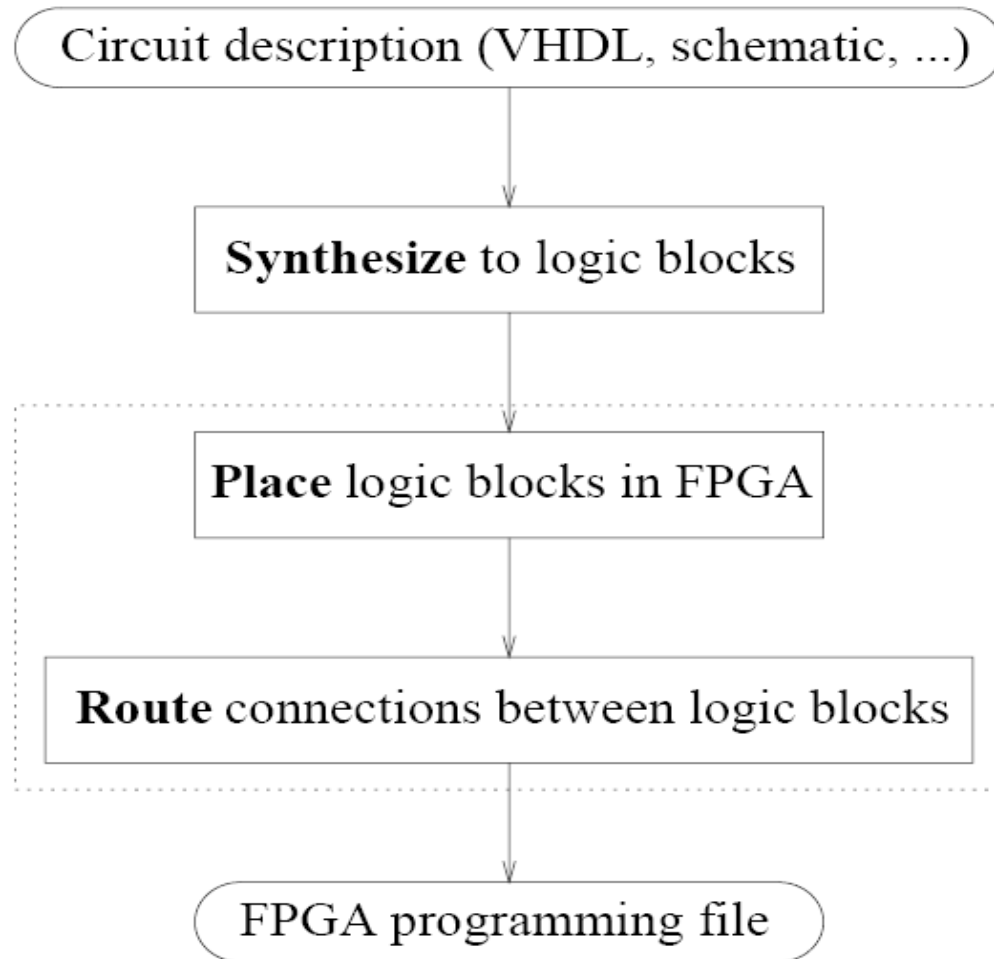
FPGA Placement

Outline

- Constraints and objectives
- Placement approaches
 - Simulated annealing-based
 - Partitioning-based
 - Analytical method
- Representative placers
 - VPR placer (SA-based) [1, 2]
 - PPFF (partitioning-based) [3]
 - Routing-architecture-aware analytical placer (analytical) [4]




FPGA CAD Flow



FPGA Placement


- Goal: Place each logic block of a netlist to a unique and legal position on an FPGA



Placement Objectives


- Minimize the required wiring
 - *wirelength-driven* placement.
- Balance the wiring density across the FPGA
 - *routability-driven* placement.
- Maximize circuit speed
 - *timing-driven* placement.

FPGA-Specific Placement Issues



- Has to deal with *fixed carrier dimensions*.
- *Channel density* in every channel cannot exceed the number of routing tracks available.
 - How to estimate the channel density?

FPGA-Specific Placement Issues



- FPGA contains routing tracks of various lengths.
 - How to estimate interconnection delay?
- Simple interconnection *delay estimation* model based on net length or fanout are not accurate for FPGA.

On Delay Estimation




Office



Home


Which is the best route from home to the office?

Longer, but
faster



Shorter, but
slower

On Delay Estimation




- central block
- 1 hop
- 2 hops
- 3 hops

Block Position	A	B
Distance	(2,1)	(5,0)
HPWL	3	5
# of hops	2	1
Delay (ps)	405	309


Close Relation with Routing

- Ideally, placement and routing (P & R) should be performed simultaneously as they depend on each other's results.
- In practice, placement is done prior to routing as simultaneous placement and routing is too complicated.




Placement Approaches

- 3 major classes of placers:
 - *Simulated annealing* based placers
 - *Min-cut (partitioning-based)* placers
 - *Analytic* placers



Simulated Annealing

- Kirkpatrick, Gelatt, and Vecchi,
“Optimization by simulated annealing”,
Science, May 1983.



Simulated Annealing Basics

- Simulated annealing (SA) is a *stochastic search technique* to search for a near optimal solution for an optimization problem.
- SA mimics the annealing process used to gradually cool molten metal to produce a high-quality crystal structure.
- SA takes an existing solution and then makes successive changes in a series of random moves.
- Each move is accepted or rejected based on an *energy function*.

Simulated Annealing Basics

- Non-zero probability for *“up-hill” moves*.
- Probability depends on
 1. magnitude of the “up-hill” movement
 2. temperature

$$Prob(S \rightarrow S') = \begin{cases} 1 & \text{if } \Delta C \leq 0 \quad /* \text{“down - hill” moves} */ \\ e^{-\frac{\Delta C}{T}} & \text{if } \Delta C > 0 \quad /* \text{“up - hill” moves} */ \end{cases}$$

- $\Delta C = cost(S') - Cost(S)$
- T : Control parameter (*temperature*)
- Annealing schedule:
 - $T = T_0, T_1, T_2, \dots$, where $T_i = r^i T_0, r < 1$.

Generic Simulated Annealing Algorithm

```
1 begin
2 Get an initial solution  $S$ ;
3 Get an initial temperature  $T > 0$ ;
4 while not yet “frozen” do
5   for  $1 \leq i \leq P$  do
6     Pick a random neighbor  $S'$  of  $S$ ;
7      $\Delta \leftarrow \text{cost}(S') - \text{cost}(S)$ ;
8     /* downhill move */
9     if  $\Delta \leq 0$  then  $S \leftarrow S'$ 
10    /* uphill move */
11    if  $\Delta > 0$  then  $S \leftarrow S'$  with probability  $e^{-\frac{\Delta}{T}}$ ;
12   $T \leftarrow rT$ ; /* reduce temperature */
13 return  $S$ 
14 end
```

Basic Ingredients for Simulated Annealing

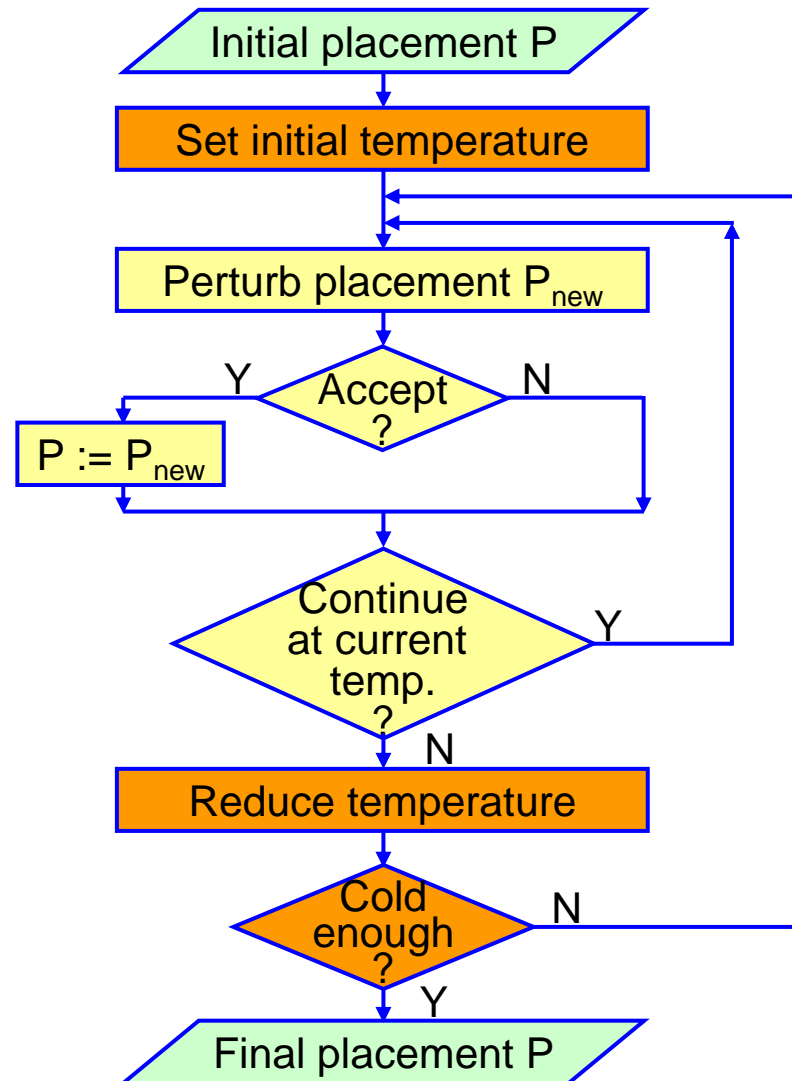
■ Analogy:

Physical system	Optimization problem
state	configuration
energy	cost function
ground state	optimal solution
quenching	iterative improvement
careful annealing	simulated annealing

■ Basic Ingredients for Simulated Annealing:


- ☐ **Solution space**
- ☐ **Neighborhood structure**
- ☐ **Cost function**
- ☐ **Annealing schedule**

Simulated Annealing Based Placement



Simulated Annealing Based Placement

- Get initial placement by assigning logic blocks of the circuit randomly to the available locations in the FPGA.
- Cost function
 - For wirelength-driven placement, a common cost function is the sum over all nets of the *half-perimeter* of their *bounding boxes*.





- Move types
 - Exchange locations of two randomly selected logic blocks
 - Move a logic block to an empty location

VPR Placer (SA-based) [1]

- VPR: versatile place and route
- A *simulated annealing-based* placer.
- *Modified temperature updating scheme*
 - Accelerated temperature decrease when move acceptance rate α is very high or low
 - $T_{new} = r T_{old}$
$$r = \begin{cases} 0.5 & \text{if } \alpha > 0.95 \\ 0.9 & \text{if } 0.8 < \alpha \leq 0.95 \\ 0.95 & \text{if } 0.15 < \alpha \leq 0.8 \\ 0.8 & \text{if } \alpha \leq 0.15 \end{cases}$$
- Fast *incremental net bounding box updating*.

VPR Placer

- Gradually reduce the scope of cell moves or exchanges (long range to short range) to *keep the move acceptance rate* close to 0.44.



VPR Placer

- *Congestion model* for non-uniform channel capacity (penalize solutions requiring more routing in the narrower channels).
- Cost function of basic mode:

$$Cost = \sum_{n=1}^{N_{nets}} q(n) \left[\frac{bb_x(n)}{C_{av,x}(n)} + \frac{bb_y(n)}{C_{av,y}(n)} \right]$$

- $bb_x(n)$: horizontal span of net n
- $bb_y(n)$: vertical span of net n
- $q(n)$: *adjustment factor* depending on #terminals of net n

Timing-Driven VPR Placer [2]

- Add a timing cost to the objective function
- Timing cost is the *weighted delays* of all connections (the weight reflects a connection's *criticality*)
 - $w(e) = (1 - \text{slack}(e)/T)^\beta$ where T is the current longest path delay, β is a constant
- *Self-normalization*: changes of timing cost and wirelength of a move are normalized by their previous values

Timing-Driven VPR Placer


- Use pre-computed *delay lookup matrix* – delay for $(\Delta x, \Delta y)$
- Run *static timing analysis* periodically to update slacks and hence criticalities

Partitioning-based Placement

- The netlist is repeatedly partitioned into two sub-circuits.
- Meanwhile, the chip area is *partitioned* into two sub-regions *recursively*.
- Each sub-circuit is assigned to a sub-region.
- The process is repeated until each sub-circuit consists of a single logic block and has a unique location on the chip area.

Partitioning-based Placement

- During partitioning, *minimize # cut nets* based on the intuition that densely connected sub-circuits should be placed closely.



PPFF Placer (partitioning-based) [3]

■ *Top-down partitioning-based* placement

- Fast (divide&conquer)

■ Delay estimation

- Empirical routing delay analysis


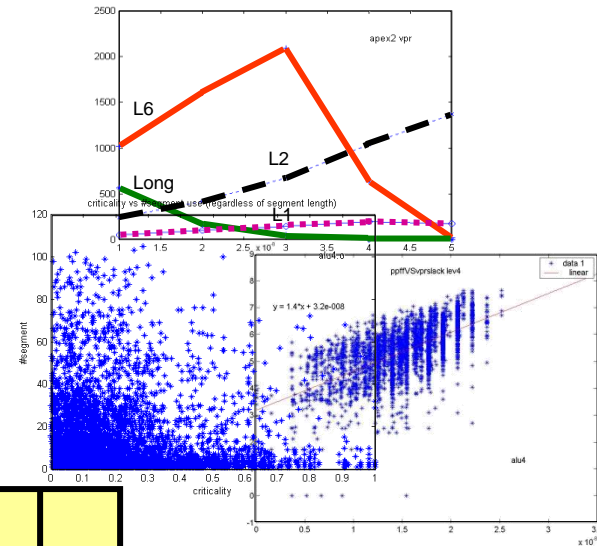

- Correlate net criticality and routing resource usage

■ Delay optimization


- Align critical connections

- Dynamically update net delays

- Net weight to represent timing criticality




Net Terminal Alignment in PPFF







- Put node X into subregion A or B?
- Alignment helps reduce # routing segments
- Y treated as anchor point
- **X → A preferred** (one segment possible)

Partitioning Order in PPFF

- Top-to-bottom and left-to-right.
- Alignment dependence "propagates"




-  *partitioned placement regions*
-  *to be partitioned placement regions*
-  *To be horizontally aligned*
-  *To be vertically aligned*


- Also proposed a more sophisticated method to determine a good partitioning order.

Alignment Details

- Higher priority for timing critical node to align with its anchor in the direction of a cut



Delay Estimation




- Wire length / bounding box **NOT** accurate
 - Same BBox, fewer segs
→ smaller delay
(e.g. A vs B)

PPFF Delay Estimation

■ Delay estimation

- *Correlate delay with distance and criticality*
- Pre-computed table based on empirical data by a router on a large no. of circuits



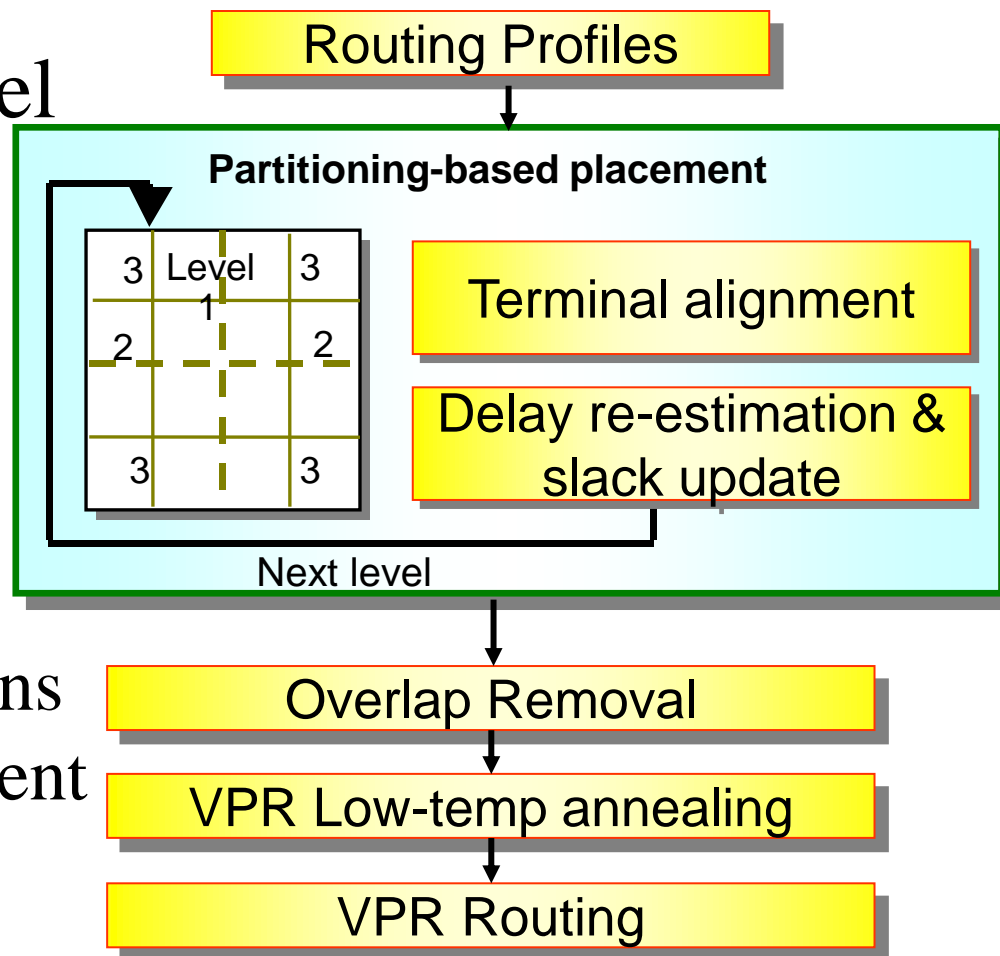
- Data shows delay as a function of net criticality and distance
- Delay is net delay after VPR routing
- Each point is the average over all nets falling in the (dist, crit) range.

Edge Weight

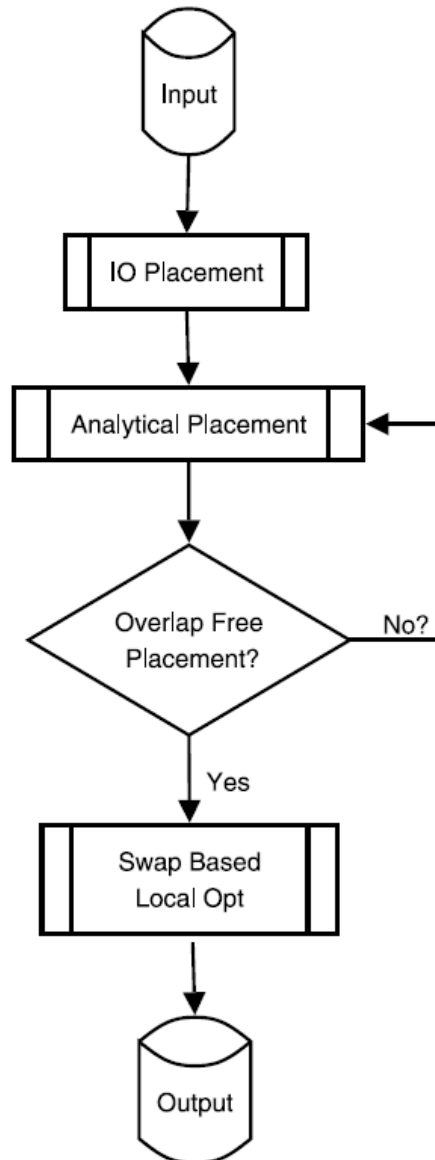
- Each multi-terminal net is decomposed into multiple 2-terminal nets (edges)
- Edge weight indicates its timing criticality
 - large edge weight will discourage partitioner to cut it
- At the end of each partitioning level
 - run static timing analysis to update slack of each edge
 - update criticality or weight of each edge based on updated slacks

PPFF Summary

- Initial analysis
- At each partitioning level
 - terminal alignment
 - delay re-estimation and slack update
- Refinement
 - Move non-critical CLB from overcrowded regions while preserving alignment
 - *Low-temperature annealing* to further optimize the placement








Analytical Placement



- Used in Xilinx's placer
- Represent placement objective and constraints as analytical functions of x_i and y_i
- Constraints are modified iteratively to move logic blocks away from highly congested regions
- After a legal placement is found, result is further optimized by local swap

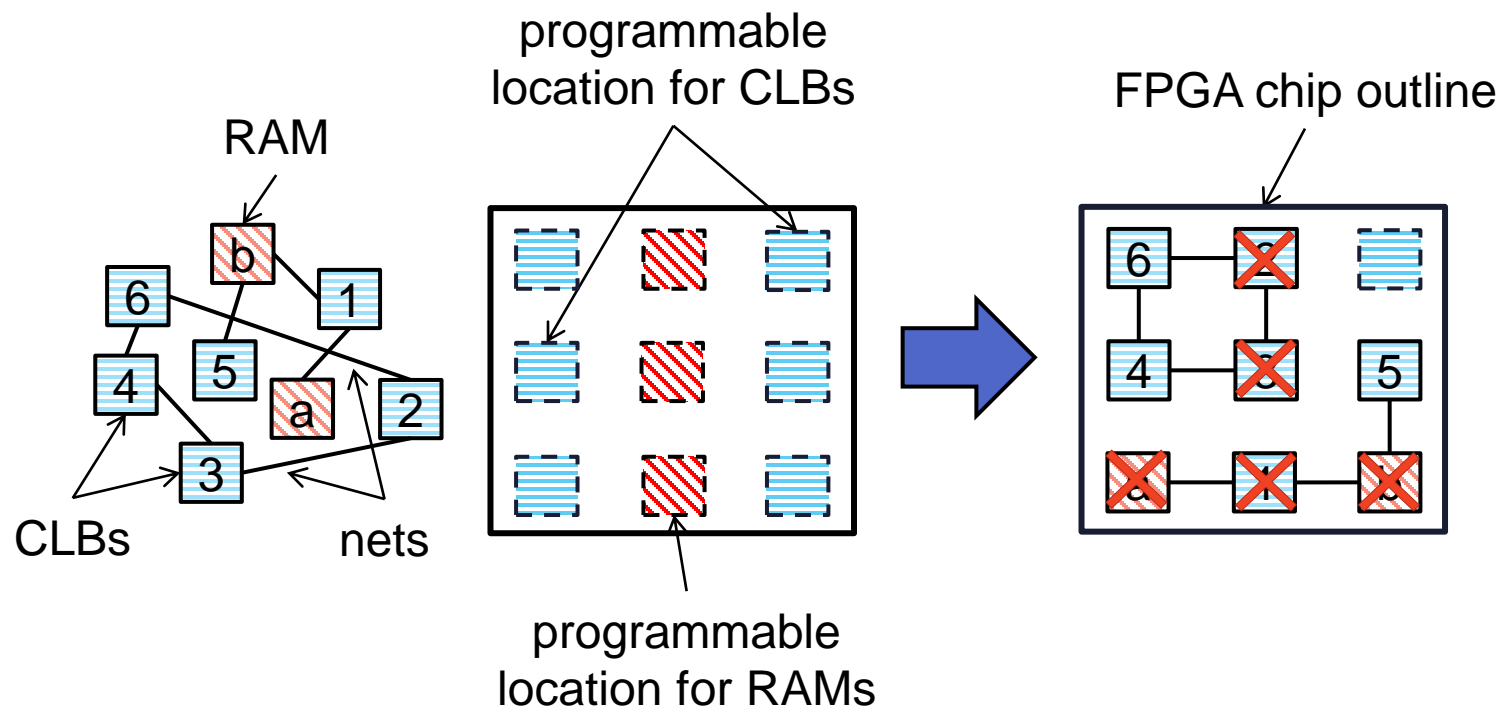
Heterogeneous FPGA

- Traditionally, FPGA consists of CLB, IO, and programmable routing arch
- Nowadays, main FPGA components:
 - Configurable logic block (**CLB** )
 - Input/output block (**IOB** )
 - Programmable routing architecture
 - Random access memory (**RAM** )
 - Digital signal processing (**DSP** )




Heterogeneous FPGA Placement

- Place each block of a netlist to a unique and legal position on an FPGA




Heterogeneous FPGA Placement

- Place each block of a netlist to a unique and legal position on an FPGA



Analytical Placement Formulation [4]

- Given the chip region and block dimensions, determine (x, y) for all movable blocks



min $W(x, y)$ // wirelength function

s.t. $D_b(x, y) \leq M_b$

D_b : density for bin b

M_b : max density for bin b

$$\text{Density} = \frac{A_{\text{block}}}{A_{\text{bin}}}$$

- Relax the constraints into the objective function (penalty)

$$\text{min } W(x, y) + \lambda \sum (\max(D_b(x, y) - M_b, 0))^2$$

- Apply **differentiable** wirelength and density models
- Use the gradient method to solve the optimization problem
- Increase λ gradually to meet density constraints

Differentiable Wirelength Model


- Log-sum-exp wirelength model [Naylor et al., 2001]

$$\gamma \sum_{e \in E} (\log \sum_{v_k \in e} \exp(x_k/\gamma) + \log \sum_{v_k \in e} \exp(-x_k/\gamma) + \log \sum_{v_k \in e} \exp(y_k/\gamma) + \log \sum_{v_k \in e} \exp(-y_k/\gamma)).$$

- Is an effective **differentiable** function for HPWL approximation and approaches exact HPWL when $\gamma \rightarrow 0$



Differentiable Density Model

- Bell-shaped density model [Kahng *et al.*, ICCAD'04]



$$D_b(\mathbf{x}, \mathbf{y}) = \sum_{v \in V} P_x(b, v) P_y(b, v)$$


$$D_b'(\mathbf{x}, \mathbf{y}) = \sum_{v \in V} c_v p_x(b, v) p_y(b, v)$$



Continuous & differentiable

Mismatch between GP and Legalization Results

- Global placement → Continuous solutions
- Legalization → Discrete and scattered legal locations
- Issue: DSP/RAM blocks get large displacement from their global placement positions



Separate Density Model

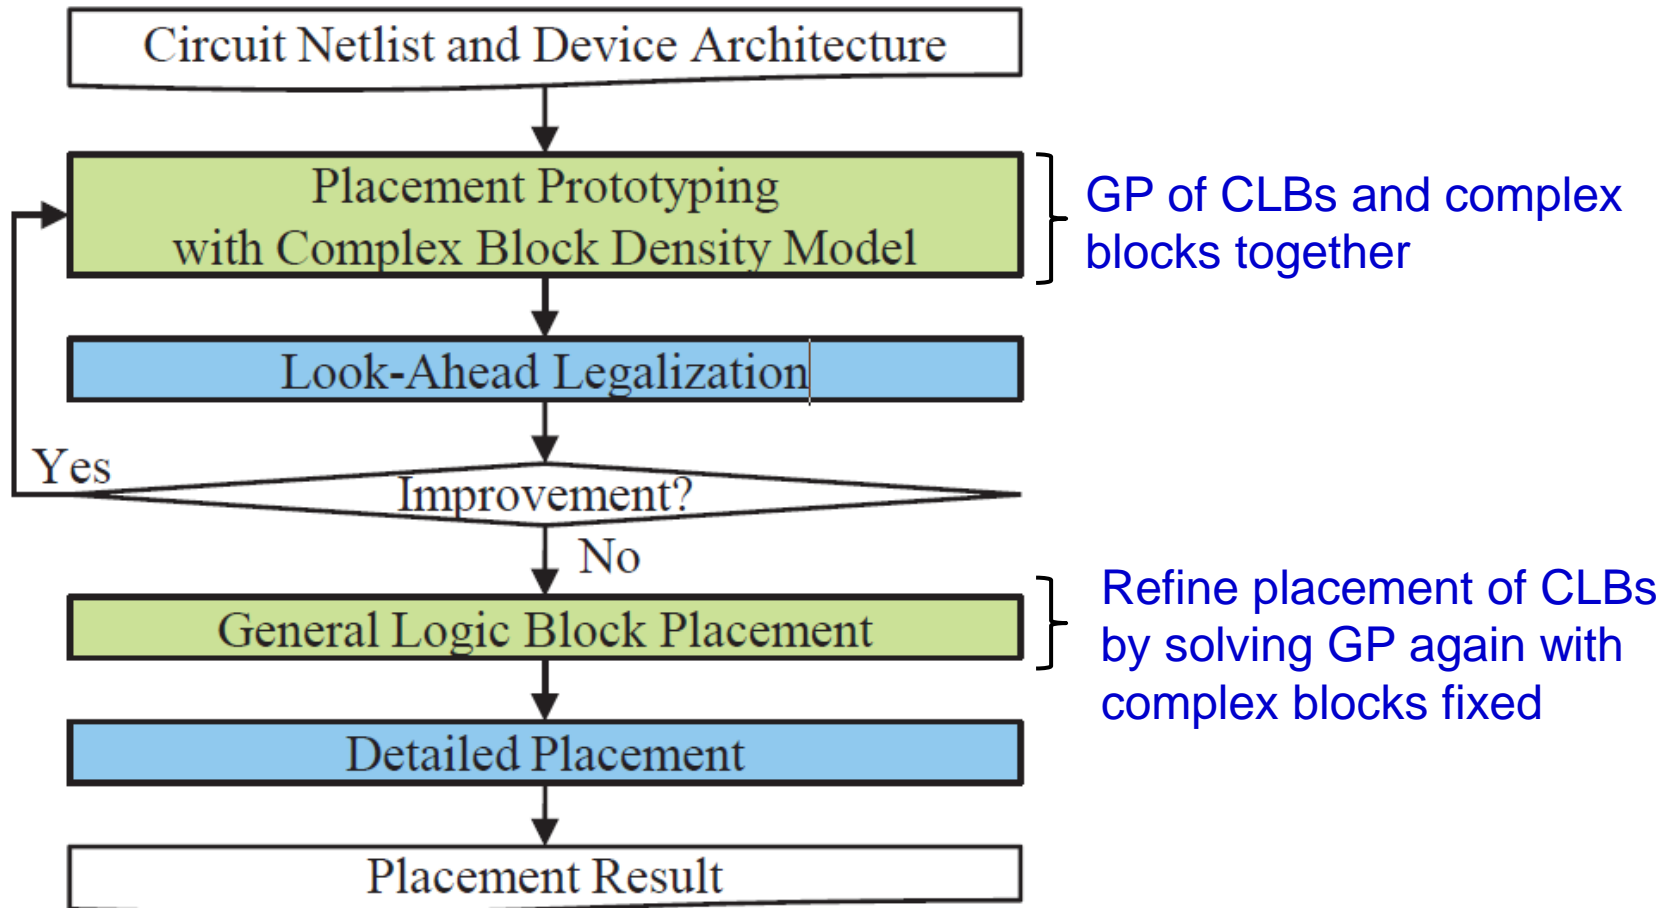
- Setup a density constraint for each type of resource
e.g. $D_b^R(\mathbf{x}, \mathbf{y}) \leq M_b^R$ where
 D_b^R : RAM density in bin b
 M_b^R : max RAM density in bin b
- Unconstrained formulation after “smoothing” and constraint relaxation with quadratic penalty method

$$\begin{aligned} \min \quad & \hat{W}(\mathbf{x}, \mathbf{y}) + \lambda \sum_b \max(\hat{D}_b(\mathbf{x}, \mathbf{y}) - M_b, 0)^2 \\ & + \lambda_R \sum_b \max(\hat{D}_b^R(\mathbf{x}, \mathbf{y}) - M_b^R, 0)^2 \\ & + \lambda_D \sum_b \max(\hat{D}_b^D(\mathbf{x}, \mathbf{y}) - M_b^D, 0)^2, \end{aligned}$$

Look-Ahead Legalization


- Give a quick forecast of legalized placement
 - Achieve more accurate wirelength estimation
 - Speedup the convergence of placement solutions
- RAMs and DSPs
 - Move blocks to the closest legal column
 - Apply bipartite matching for blocks and empty slots on each column
- General logic blocks
 - Apply the Tetris algorithm [Hill, 2002] which greedily packs blocks from the left to the right

Overall Flow



Detailed Placement

- Apply window-based cell matching and cell swapping to further reduce wirelength

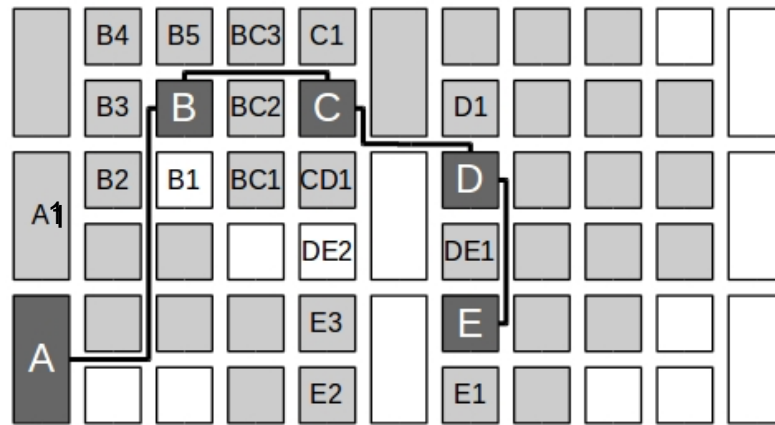


Timing-Driven Detailed Placement

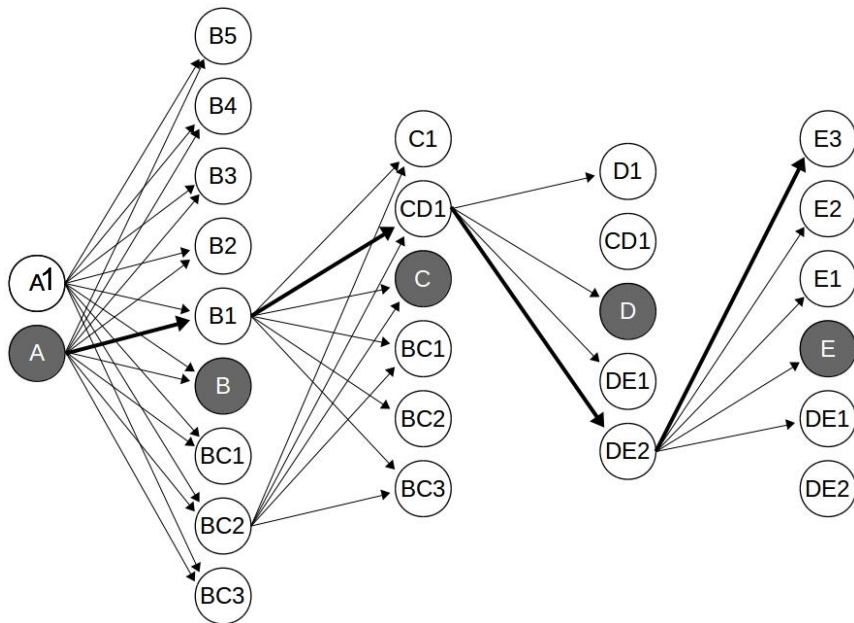
- Timing-driven placement approaches can be net-based or path-based
 - Net-based approach assigns higher net weights for more critical nets, then minimizes the weighted wirelength (e.g. timing-driven VPR)
 - Path-based approach tries to optimize critical path delay directly

Timing-Driven Detailed Placement [5]


- [5] proposed a detailed placement algorithm to optimize critical paths one at a time
- Optimization of a critical path is modeled as a shortest path problem
- Each element on the selected path may move within its neighborhood s.t. no move/swap should increase the worst negative slack in the circuit



(a) A critical path A-B-C-D-E



(b) Formulate a shortest path problem
(not all edges are shown)



(c) Change placement accordingly

References

- [1] “VPR: A New Packing, Placement and Routing Tool for FPGA Research,” in *FPL’97*
- [2] “Timing-Driven Placement for FPGAs”, in *FPGA’00*
- [3] “Fast Timing-Driven Partitioning-based Placement for Island Style FPGAs”, in *DAC’03*
- [4] “Routing Architecture-Aware Analytical Placement for Heterogeneous FPGAs”, in *DAC’15*.
- [5] “An Effective Timing-Driven Detailed Placement Algorithm for FPGAs” in *ISPD’17*.