



CS5120 VLSI System Design, Spring 2024

# DNN Mapping Part 2: Tiling for Hardware Structure

黃稚存

Chih-Tsun Huang

[cthuang@cs.nthu.edu.tw](mailto:cthuang@cs.nthu.edu.tw)



國立清華大學  
NATIONAL TSING HUA UNIVERSITY

資訊工程學系  
Computer Science

Lecture 07

- ◎ 本課程之內容 (包括但不限於教材、影片、圖片、檔案資料等)，僅供修課學生個人合理使用，非經授課教師同意，不得以任何形式轉載、重製、散布、公開播送、出版或發行本影片內容 (例如將課程內容放置公開平台上，如 Facebook, Instagram, YouTube, Twitter, Google Drive, Dropbox 等等)。如有侵權行為，需自負法律責任。




# Outline

- ⦿ Hardware Resource Constraints
- ⦿ Mapping Space
- ⦿ Tuning of DNN Mapping



# Recap: Matrix Multiplication for Fully-Connected Layer


- Consider  $A * B = C$ 
  - Can be generalized to convolution





# Recap: Loop Nest for Matrix Multiplication

- Consider  $A * B = C$

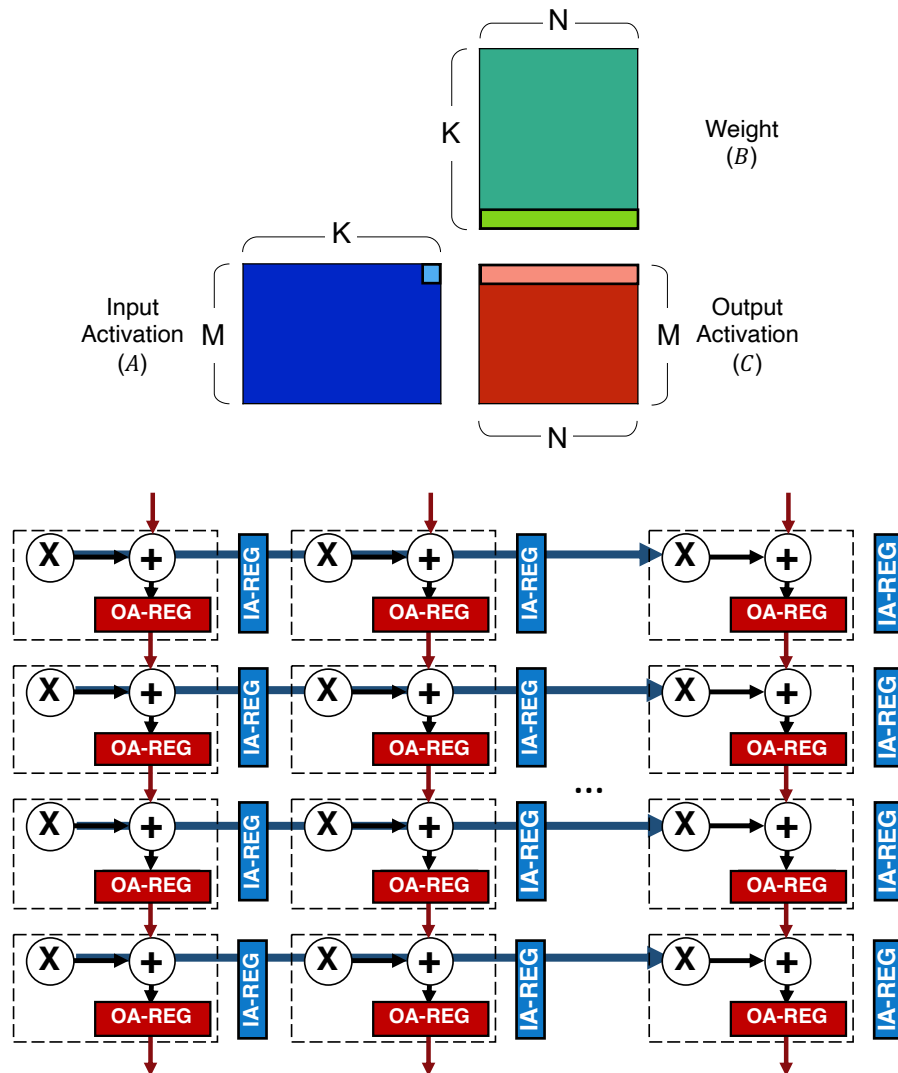


```
for (m=0; m<M; m++) {  
    for (n=0; n<N; n++) {  
        OA[n,m] = 0;  
        for (k=0; k<K; k++) {  
            OA[n,m] += IA[m, k] * W[k, n];  
        }  
        OA[n,m] = Activation(OA[n,m]);  
    }  
}
```

For each output activation

Reduction

# Recap: Mixed Datapath Optimization: TPU

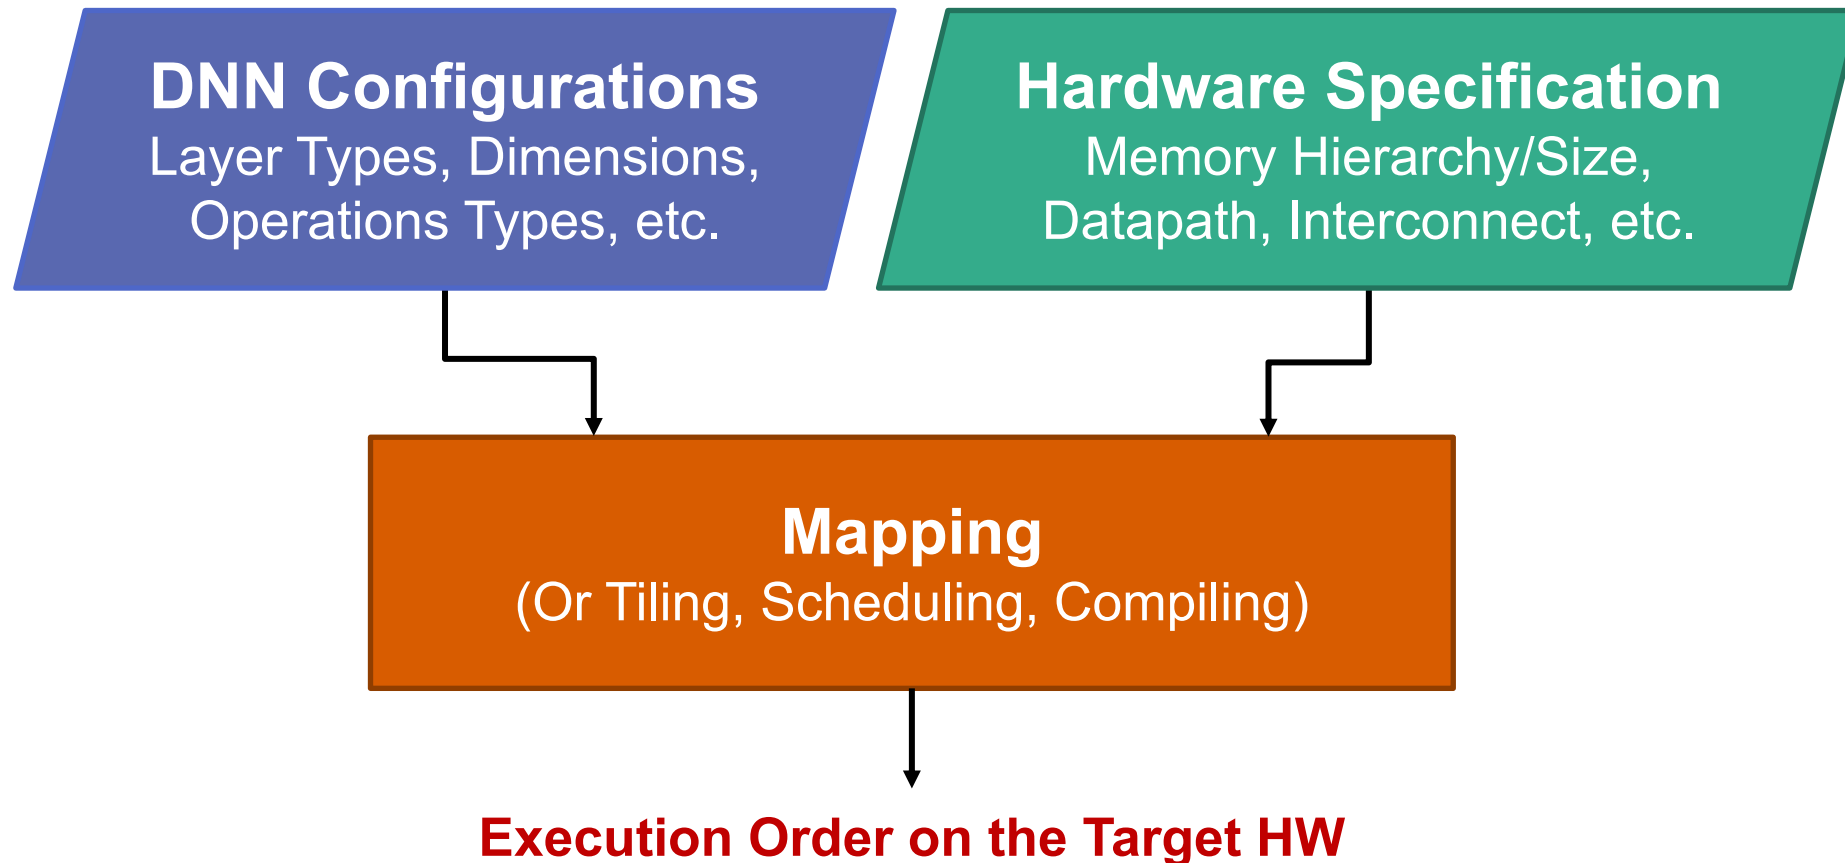


```
for (m=0; m<M; m++) {  
  parallel_for (n=0; n<N; n++) {  
    OA[n,m] = 0;  
    parallel_for (k=0; k<K; k++) {  
      OA[n,m] += IA[m,k] * W[k,n];  
    }  
    OA[n,m] = Activation(OA[n,m]);  
  }  
}
```

- Systolic accumulation
- Systolic multicast
- Area vs. scalability
- Latency vs. pipeline throughput



# DNN Mapping Problem



- ⊙ Mapping objective: efficient for the performance (latency) and/or energy, etc.




# Hardware Resource Constraints




# Constraint 1: Systolic Array Size $< K$ and/or $N$

- DNN Configuration:  
Matrix Multiplication



- Hardware Specification:  
Systolic array size:  $2 \times 2$



- Notation:  $K?/N? \rightarrow$  loop bounds

$$N = N_1 \times N_0$$

$$K = K_1 \times K_0$$

```
for (m=0; m<M; m++) {  
  for (n1=0; n1<N1; n1++) {  
    OA[n1*N0:(n1+1)*N0,m] = 0; } } } } }  
    for (k1=0; k1<K1; k1++) {  
      parallel_for (n0=0; n0<N0; n0++) {  
        parallel_for (k0=0; k0<K0; k0++) {  
          OA[n1*N0+n0,m] += IA[m,k1*K0+k0]  
            * W[k1*K0+k0,n1*N0+n0];  
        }  
      }  
    }  
  }  
}
```

Temporal  
Tiling

Spatial  
Tiling

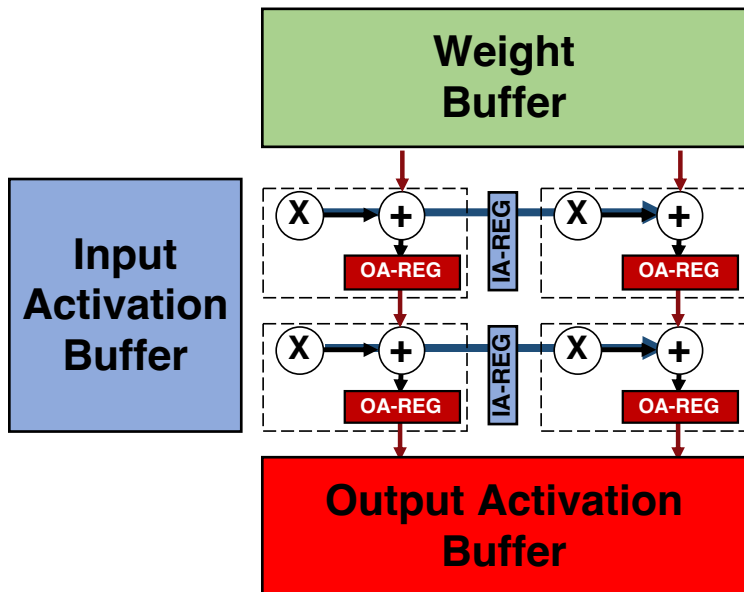

在空間上展開

}}}}}}



# Constraint 2: Weight Buffer Size $< K * N$

## (1/3)



### Hardware Specification:

- ◆ Systolic array size: 2x2
- ◆ Explicit data movement (or data orchestration)

```
for (m=0; m<M; m++) {
```

```
    for (n1=0; n1<N1; n1++) {
```

```
        OA[n1*N0:(n1+1)*N0,m] = 0;
```

```
        for (k1=0; k1<K1; k1++) {
```

```
            parallel_for (n0=0; n0<N0; n0++) {
```

```
                parallel_for (k0=0; k0<K0; k0++) {
```

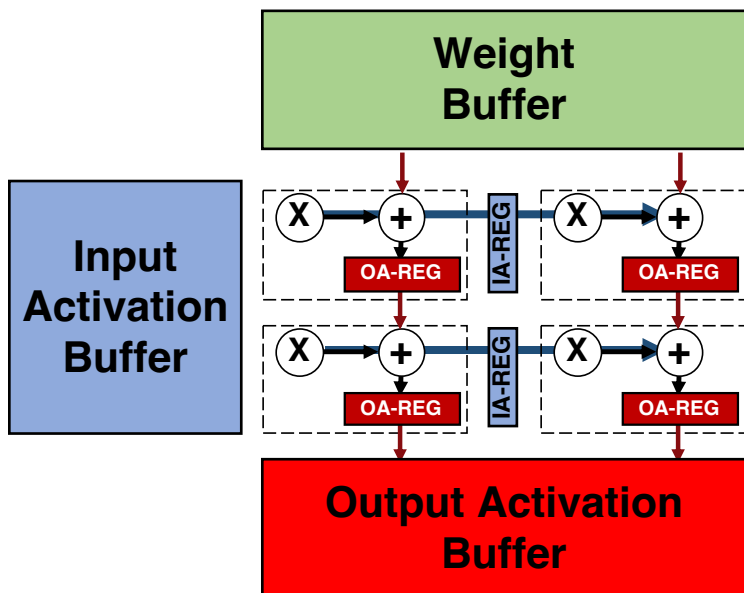

```
                    OA[n1*N0+n0,m] += IA[m,k1*K0+k0]  
                        * W[k1*K0+k0,n1*N0+n0];
```

```
                }  
            }  
        }  
    }
```

```
}
```

# Constraint 2: Weight Buffer Size < K \* N

## (2/3)



### Hardware Specification:

- ◆ Systolic array size: 2x2
- ◆ Explicit data movement (or data orchestration)
- ◆ Weight, input activation, output activation **buffer sizes**

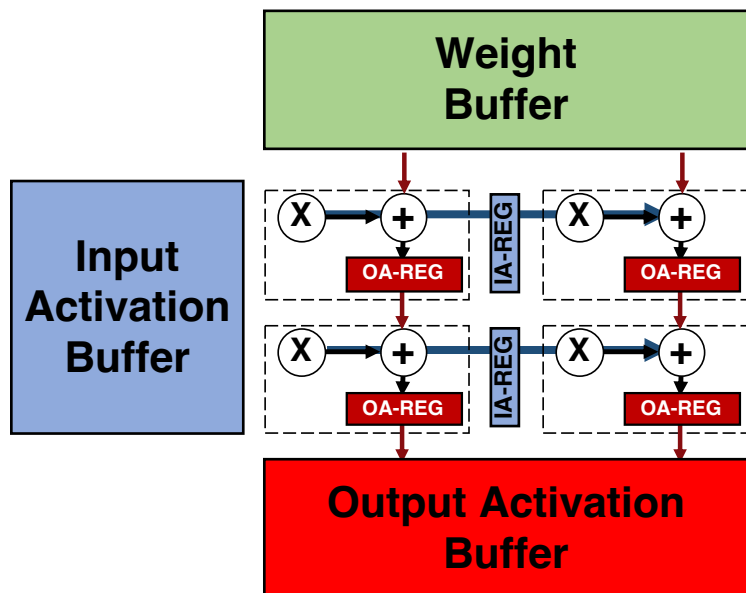

*move data to buffer*

```

for (m=0; m<M; m++) {
    mvin(W[0:K,0:N]);           // W buffer >= K*N 整个weight matrix
    mvin(IA[m:m+1,0:K]);        // IA buffer >= 1*K 一行input activation
    for (n1=0; n1<N1; n1++) {
        OA[n1*N0:(n1+1)*N0,m] = 0;
        for (k1=0; k1<K1; k1++) {
            parallel_for (n0=0; n0<N0; n0++) {
                parallel_for (k0=0; k0<K0; k0++) {
                    OA[n1*N0+n0,m] += IA[m,k1*K0+k0]
                                   * W[k1*K0+k0,n1*N0+n0];
                }
            }
        }
    }
    mvout(OA[0:N,m:m+1]);       // OA buffer >= 1*N
}
    
```

# Constraint 2: Weight Buffer Size < K \* N

## (3/3)



```

for (m=0; m<M; m++) {
    // IA buffer: 1*K
    mvin(IA[m:m+1,0:K]);
    for (n2=0; n2<N2; n2++) {
        // W buffer: N1*N0*K < K*N
        mvin(W[0:K,n2*N1*N0:(n2+1)*N1*N0]);
        OA[n2*N1*N0:(n2+1)*N1*N0,m:m+1]=0;
        for (n1=0; n1<N1; n1++) {
            for (k1=0; k1<K1; k1++) {
                parallel_for (n0=0; n0<N0; n0++) {
                    parallel_for (k0=0; k0<K0; k0++) {
                        OA[n2*N1*N0+n1*N0+n0,m]
                            += IA[m,k1*K1+k0]
                               * W[k1*K0+k0,n2*N1*N0+n1*N0+n0];
                    }
                }
            }
        }
        mvout(OA[n2*N1*N0:(n2+1)*N1*N0,m:m+1]);
    }
}

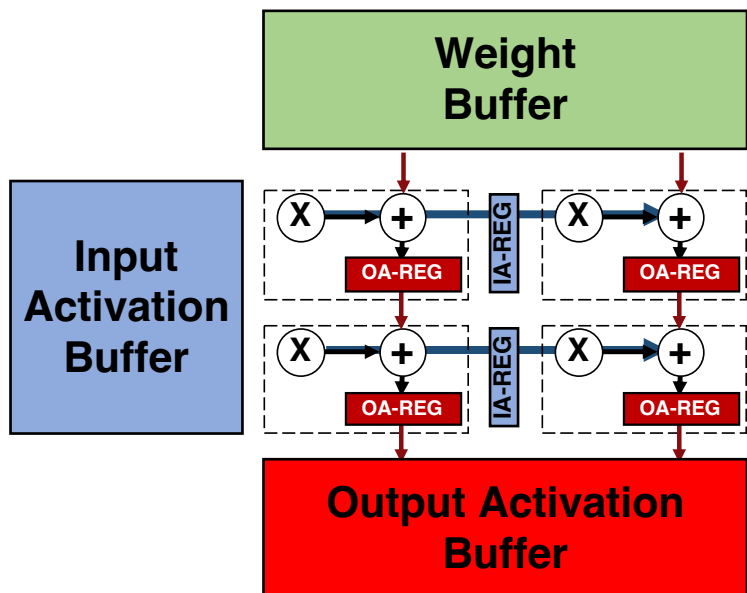

```

$N = N_2 \times N_1 \times N_0$

放不下整個 weight matrix

$N_1 \times N_0$

# Constraint 3: Input Buffer Size $> 1 * K$



for (m2=0; m2<M2; m2++) { *input activation buffer 放大 M1 倍*

// IA buffer:  $M1 * K > 1 * K$

**mvin**(IA[m2\*M1:(m2+1)\*M1,0:K]);

for (n2=0; n2<N2; n2++) {

// W buffer:  $N1 * N0 * K < K * N$

**mvin**(W[0:K,n2\*N1\*N0:(n2+1)\*N1\*N0]);

**OA**[n2\*N1\*N0:(n2+1)\*N1\*N0,m2\*M1:(m2+1)\*M1]=0;

for (m1=0; m1<M1; m1++) { *每一輪做 M1 \* K 個*

for (n1=0; n1<N1; n1++) {

for (k1=0; k1<K1; k1++) {

parallel\_for (n0=0; n0<N0; n0++) {

parallel\_for (k0=0; k0<K0; k0++) {

OA[n2\*N1\*N0+n1\*N0+n0,m2\*M1+m1]

+= IA[m2\*M1+m1,k1\*K1+k0]


\* W [k1\*K0+k0,n2\*N1\*N0+n1\*N0+n0];

}}}}}

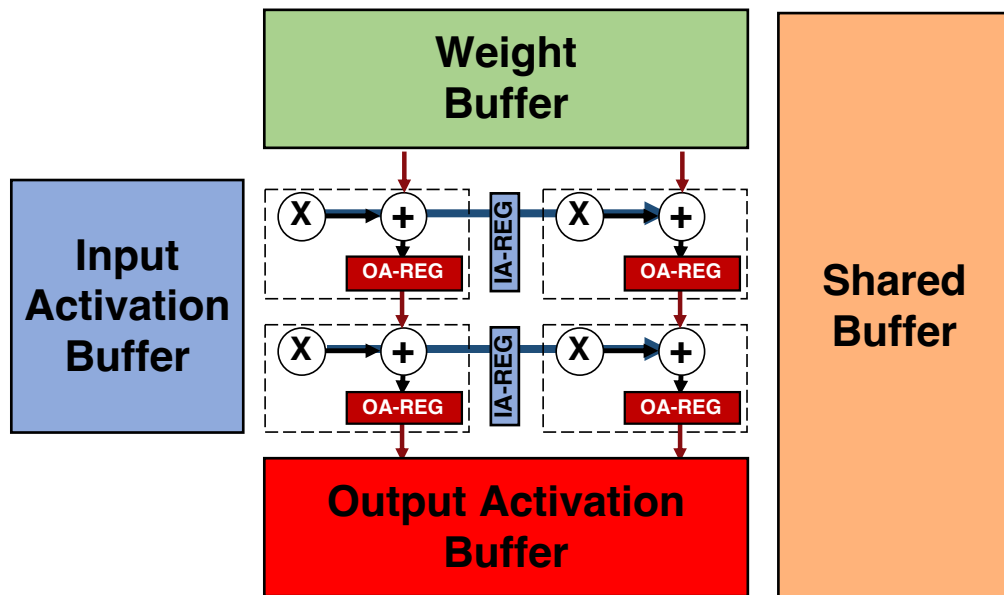
**mvout**(OA[n2\*N1\*N0:(n2+1)\*N1\*N0, m2\*M1:(m2+1)\*M1]);

}}

# Constraint 4: Adding Another Shared Buffer



*clusters grouping*



```
for (m3=0; m3<M3; m3++) {
  for (n3=0; n3<N3; n3++) {
    // Shared buffer blocking
```

```
for (m2=0; m2<M2; m2++) {
  // IA buffer stores: M1*K
  mvin(IA[...:...]);
  for (n2=0; n2<N2; n2++) {
    // W buffer stores: N1*N0*K
    mvin(W[...:...]);
    OA[...:...]=0;
    for (m1=0; m1<M1; m1++) {
      for (n1=0; n1<N1; n1++) {
        for (k1=0; k1<K1; k1++) {
          ...
        }
      }
    }
    mvout(OA[...:...]);
  }
}
```



# Mapping Space

- » Loop ordering
- » Loop bound
- » Spatial choice

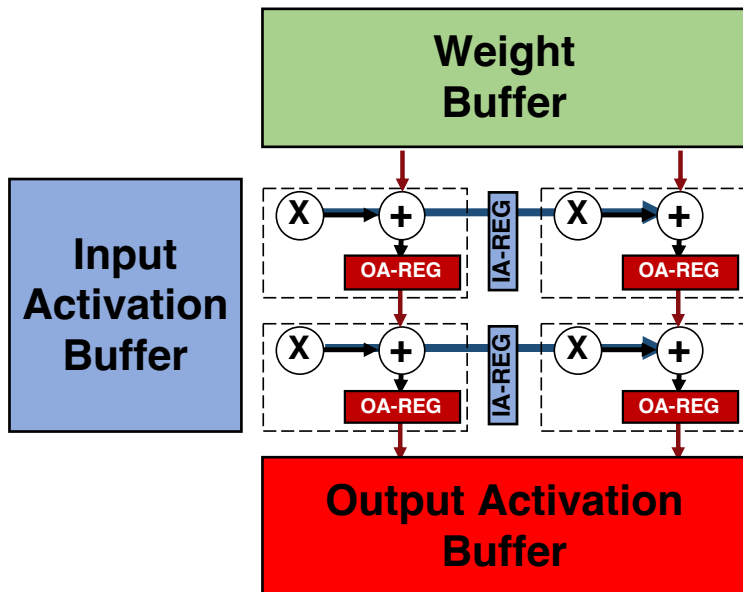



# Mapping Dimensions

- ◎ Loop ordering:
  - ◆ Which index goes to the inner/outer loop?
  
- ◎ Loop bounds:
  - ◆ What are the loop bounds (i.e.,  $N?$ / $K?$ / $M?$ ) for each loop?
  
- ◎ Spatial choice:
  - ◆ Which loop should be spatial/temporal?
    - Data/Model Parallelism



# Mapping Problem Example

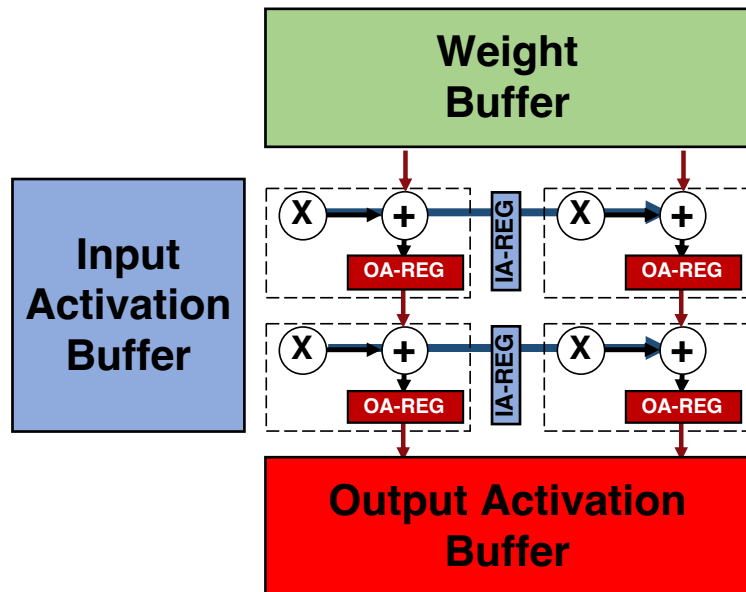



```

for (m2=0; m2<M2; m2++) {
    // IA buffer: M1*K > 1*K
    mvin(IA[m2*M1:(m2+1)*M1,0:K]);
    for (n2=0; n2<N2; n2++) {
        // W buffer: N1*N0*K < K*N
        mvin(W[0:K,n2*N1*N0:(n2+1)*N1*N0]);
        OA[n2*N1*N0:(n2+1)*N1*N0,m2*M1:(m2+1)*M1]=0;
        for (m1=0; m1<M1; m1++) {
            for (n1=0; n1<N1; n1++) {
                for (k1=0; k1<K1; k1++) {
                    parallel_for (n0=0; n0<N0; n0++) {
                        parallel_for (k0=0; k0<K0; k0++) {
                            OA[n2*N1*N0+n1*N0+n0,m2*M1+m1]
                                += IA[m2*M1+m1,k1*K1+k0]
                                    * W [k1*K0+k0,n2*N1*N0+n1*N0+n0];
                        }
                    }
                }
            }
        }
        mvout(OA[n2*N1*N0:(n2+1)*N1*N0,m2*M1:(m2+1)*M1]);
    }
}

```

# Loop Bound



```
for (m2=0; m2<M2; m2++) {
    // IA buffer: M1*K > 1*K
    mvin(IA[m2*M1:(m2+1)*M1,0:K]);
    for (n2=0; n2<N2; n2++) {
```

$$N = N_2 \times N_1 \times N_0$$

$$M = M_1 \times M_2$$

$$K = K_1 \times K_0$$

```
    // W buffer: N1*N0*K < K*N
```

```
    mvin(W[0:K,n2*N1*N0:(n2+1)*N1*N0]);
```

```
    OA[n2*N1*N0:(n2+1)*N1*N0,m2*M1:(m2+1)*M1]=0;
```

```
    for (m1=0; m1<M1; m1++) {
```

```
        for (n1=0; n1<N1; n1++) {
```

```
            for (k1=0; k1<K1; k1++) {
```

```
                parallel_for (n0=0; n0<N0; n0++) {
```

```
                    parallel_for (k0=0; k0<K0; k0++) {
```

```
                        OA[n2*N1*N0+n1*N0+n0,m2*M1+m1]
```

```
                        += IA[m2*M1+m1,k1*K1+k0]
```



```
                        * W [k1*K0+k0,n2*N1*N0+n1*N0+n0];
```

```
                    }
                }
            }
        }
    }
```

```
compute_matmul();
```

```
    mvout(OA[n2*N1*N0:(n2+1)*N1*N0,m2*M1:(m2+1)*M1]);
```

# Loop Ordering (1/2)



```

for (m2=0; m2<M2; m2++) {
    // IA buffer: M1*K > 1*K
    mvin(IA[m2*M1:(m2+1)*M1,0:K]);
    for (n2=0; n2<N2; n2++) {
        // W buffer: N1*N0*K < K*N
        mvin(W[0:K,n2*N1*N0:(n2+1)*N1*N0]);
        compute_matmul(*W, *IA, *OA,
                        N2, N1, N0,
                        M2, M1,
                        K1, K0,
                        m2, n2);
        mvout(OA[n2*N1*N0:(n2+1)*N1*N0,
                 m2*M1:(m2+1)*M1]);
    }
}
    
```

# Loop Ordering (2/2)

$$N = N_2 \times N_1 \times N_0$$
$$M = M_1 \times M_2$$
$$K = K_1 \times K_0$$

Option 1: Loops  $m_2 \rightarrow n_2$

Option 2: Loops  $n_2 \rightarrow m_2$

```
for (m2=0; m2<M2; m2++) {  
  // IA buffer:  $M_1 \times K$   $M_2 \times M_1 \times K = M \times K$   
  mvin(IA[m2*M1:(m2+1)*M1,0:K]);  
  for (n2=0; n2<N2; n2++) {  
    // W buffer:  $N_1 \times N_0 \times K$   $M_2 \times N_2 \times N_1 \times N_0 \times K$   
    mvin(W[0:K,n2*N1*N0:(n2+1)*N1*N0]);  
    compute_matmul(*W, *IA, *OA,  
      ...  
      m2, n2);  
    mvout(OA[n2*N1*N0:(n2+1)*N1*N0,  
      m2*M1:(m2+1)*M1]);  
  }  
}
```

buffer 內所有資料的移動量

```
for (n2=0; n2<N2; n2++) {  
  // W buffer:  $N_1 \times N_0 \times K$   $N_2 \times N_1 \times N_0 \times K = N \times K$   
  mvin(W[0:K,n2*N1*N0:(n2+1)*N1*N0]);  
  for (m2=0; m2<M2; m2++) {  
    // IA buffer:  $M_1 \times K$   $N_2 \times M_2 \times M_1 \times K$   
    mvin(IA[m2*M1:(m2+1)*M1,0:K]);  
    compute_matmul(*W, *IA, *OA,  $N_2 \times M \times K$   
      ...  
      m2, n2);  
    mvout(OA[n2*N1*N0:(n2+1)*N1*N0,  
      m2*M1:(m2+1)*M1]);  
  }  
}
```

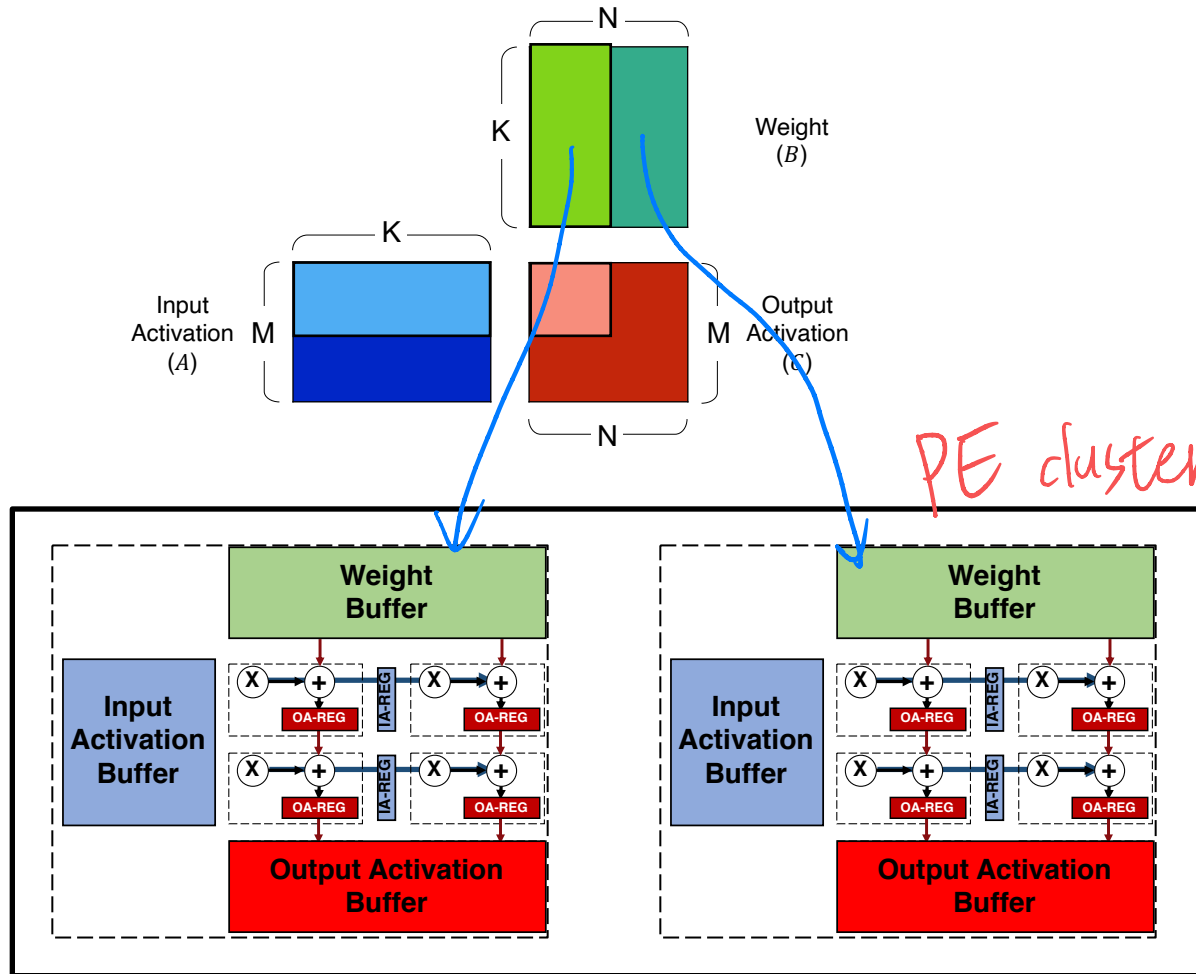
IA Movement:  $M \times K$

W Movement:  $M_2 \times N \times K$

IA Movement:  $N_2 \times M \times K$

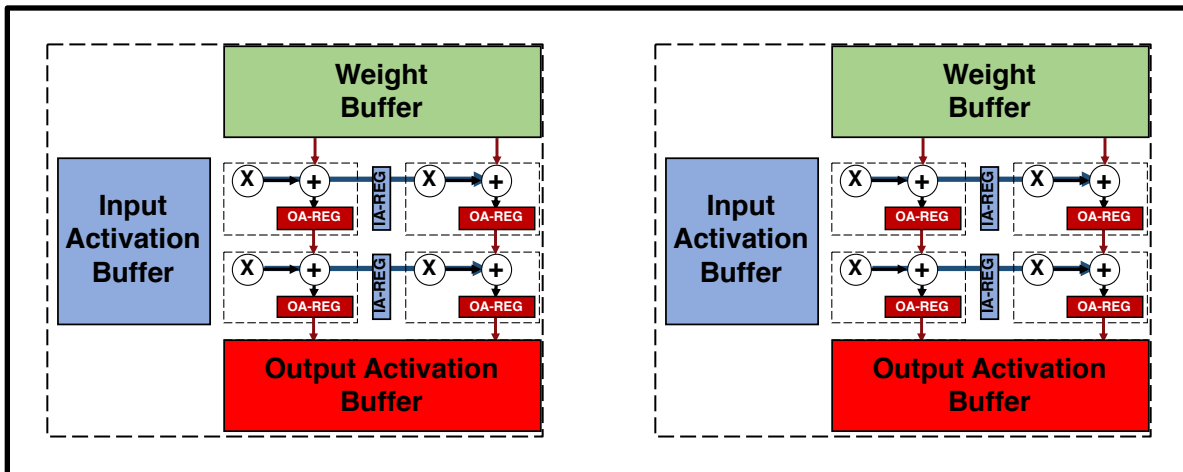

W Movement:  $N \times K$

# Spatial Choice: Model Parallelism



```
parallel_for (n2=0; n2<N2; n2++) {
    // Model Parallelism
    mvin(W[0:K, n2*N1*N0:(n2+1)*N1*N0]);
    for (m2=0; m2<M2; m2++) {
        mvin(IA[m2*M1:(m2+1)*M1, 0:K]);
        compute_matmul(*W, *IA, *OA,
            ...
            m2, n2);
        mvout(OA[n2*N1*N0:(n2+1)*N1*N0,
            m2*M1:(m2+1)*M1]);
    }
}
```

# Spatial Choice: Data Parallelism



```
parallel_for (m2=0; m2<M2; m2++) {
    // Data Parallelism
    mvin(IA[m2*M1:(m2+1)*M1,0:K]);
    for (n2=0; n2<N2; n2++) {
        mvin(W[0:K,n2*N1*N0:(n2+1)*N1*N0]);
        compute_matmul(*W, *IA, *OA,
            ...
            m2, n2);
        mvout(OA[n2*N1*N0:(n2+1)*N1*N0,
            m2*M1:(m2+1)*M1]);
    }
}
```

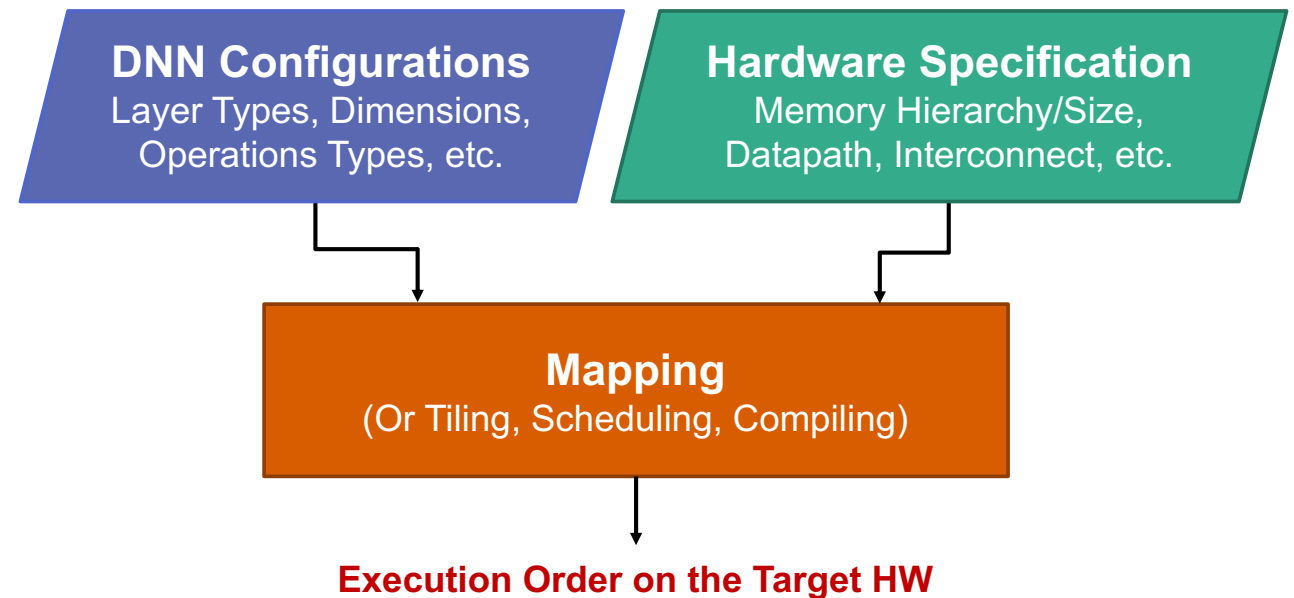


# Tuning of DNN Mapping



# Mapping Dimensions

- DNN mapping problem → an **optimization problem**
- Given:
  - ◆ DNN dimensions (N, H, W, C, R, S, K, stride, padding)
  - ◆ Hardware specifications (dataflow, memory hierarchy)
- Objective:
  - ◆ An optimal loop nest that minimizes latency and/or energy
    - Both temporal and spatial execution order
- Approaches:
  - ◆ Exhaustive search
  - ◆ Random search
  - ◆ Learning-based algorithms

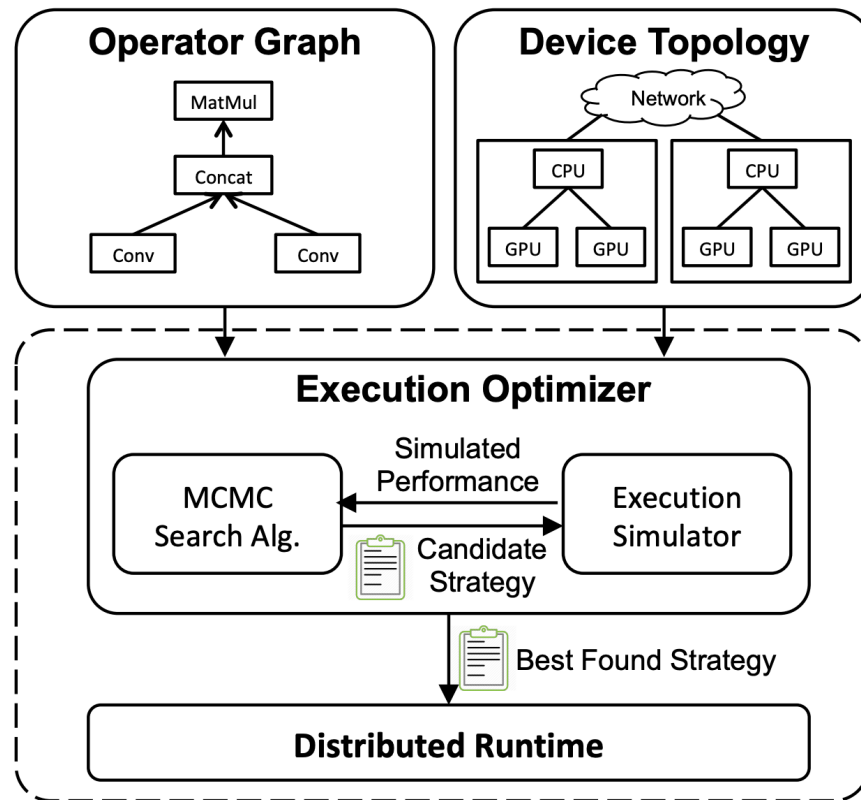






# Example: FlexFlow, SysML'2018

- “The optimizer uses a MCMC (Markov Chain Monte Carlo) search algorithm to explore the space of possible parallelization strategies and iteratively proposes candidate strategies that are evaluated by an execution simulator.”

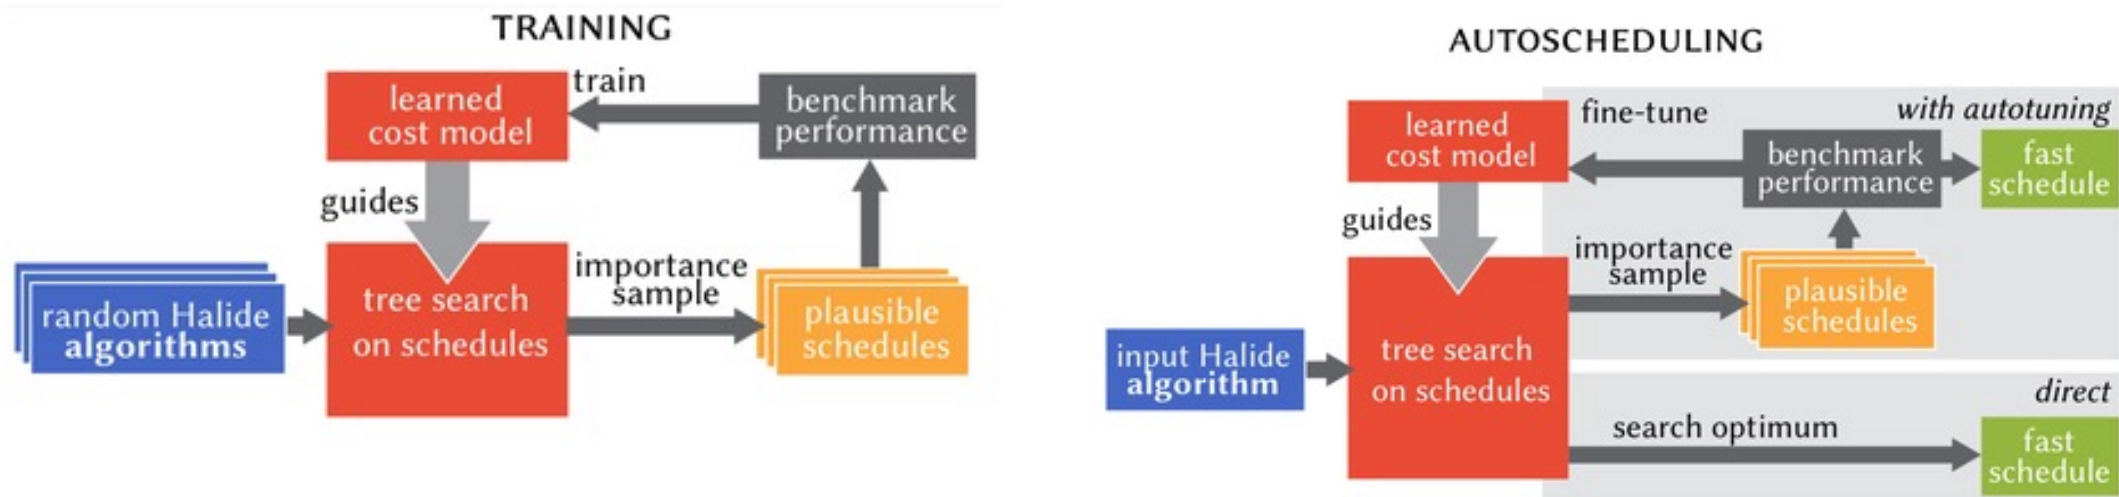


Beyond Data and Model Parallelism for Deep Neural Networks, SysML 2018



# Example: Halide, SIGGRAPH'2019

- “We generate schedules for Halide programs using tree search over the space of schedules guided by a learned cost model and optional autotuning. The cost model is trained by benchmarking thousands of randomly-generated Halide programs and schedules. The resulting code significantly outperforms prior work and human experts.”

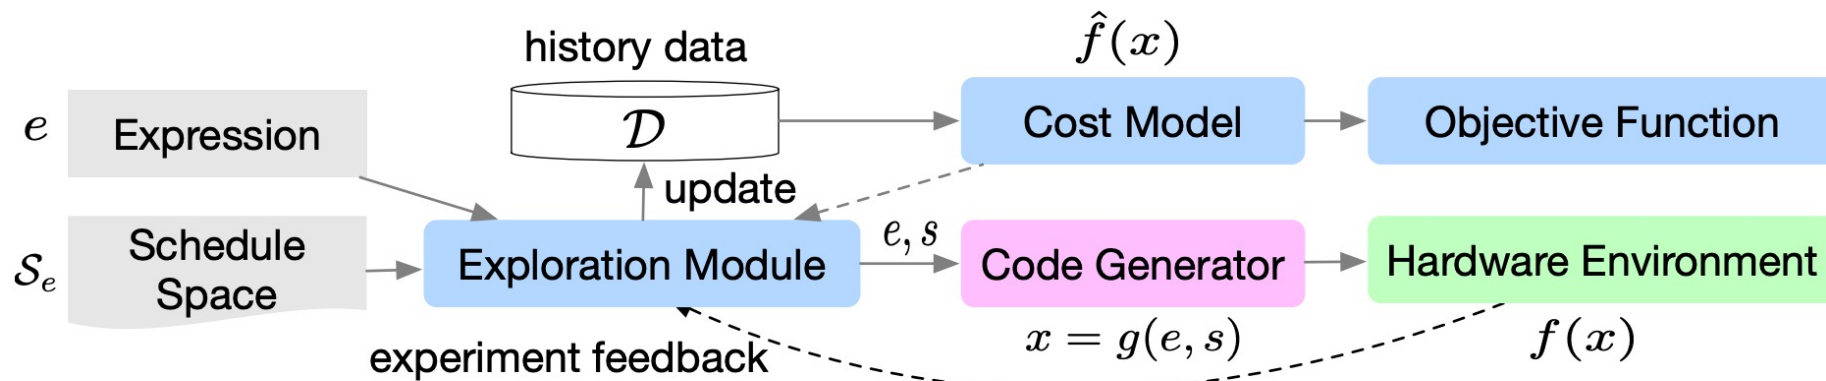
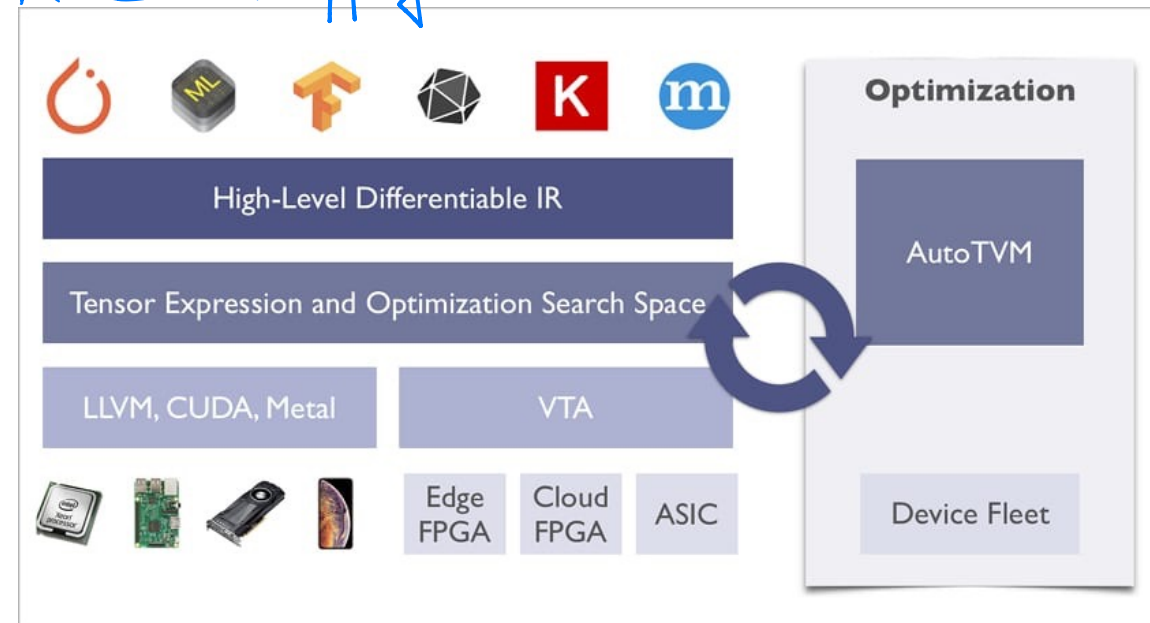


Learning to Optimize Halide with Tree Search and Random Programs, *SIGGRAPH 2019*

# Example: TVM, NeurIPS'2018

auto tuning, 找出一個最適合的 mapping 方式

- Learn domain-specific statistical cost models to guide the search of tensor operator implementations over billions of possible program variants
- Further accelerate the search using effective model transfer across workloads



Learning to Optimize Tensor Programs, *NeurIPS'2018*



# Summary

- ⊙ Temporal and Spatial Mapping based on hardware constraints
  - ◆ Memory hierarchy
  - ◆ Parallelism
- ⊙ Tile the loops to improve reuse and parallelism
  - ◆ Loop ordering
  - ◆ Loop bounds
  - ◆ Spatial choices
- ⊙ Navigate the large mapping space
  - ◆ Finding an optimized solution
  - ◆ Getting to the solution fast enough