

CS5120 VLSI System Design, Spring 2024

Data Reuse and Dataflow

黃稚存

Chih-Tsun Huang

cthuang@cs.nthu.edu.tw



國立清華大學
NATIONAL TSING HUA UNIVERSITY

資訊工程學系
Computer Science

Lecture 05



聲明

- ◎ 本課程之內容 (包括但不限於教材、影片、圖片、檔案資料等)，僅供修課學生個人合理使用，非經授課教師同意，不得以任何形式轉載、重製、散布、公開播送、出版或發行本影片內容 (例如將課程內容放置公開平台上，如 Facebook, Instagram, YouTube, Twitter, Google Drive, Dropbox 等等)。如有侵權行為，需自負法律責任。



Outline

- ⦿ Matrix Multiplication in DNN
- ⦿ Tiling
- ⦿ Data Locality and Reuse
- ⦿ Data Reuse in DNN
 - ◆ Temporal reuse
 - ◆ Spatial reuse
 - ◆ Reducing reuse distance
- ⦿ Dataflow
- ⦿ Tiled Loop Nest for Dataflow


Reference:

V. Sze, Y.-H. Chen, T.-J. Yang and J. S. Emer, "
Efficient Processing of Deep Neural Networks -- Synthesis Lectures on Computer Architecture, "
Morgan&Calypool Publishers, 2020.



Recap: Matrix Multiplication in DNN

Recap: Fully-Connected Layer



H = 1
W = 1
R = 1
S = 1
P = 1
Q = 1


U (Stride) = 1

Pad (Padding) = 0


C: # of Input Channels
M: # of Output Channels
N: Batch size



Alternative: Mapping FC Layer to Matrix Multiplication



Single Input Activation Map
Batch Size = 1

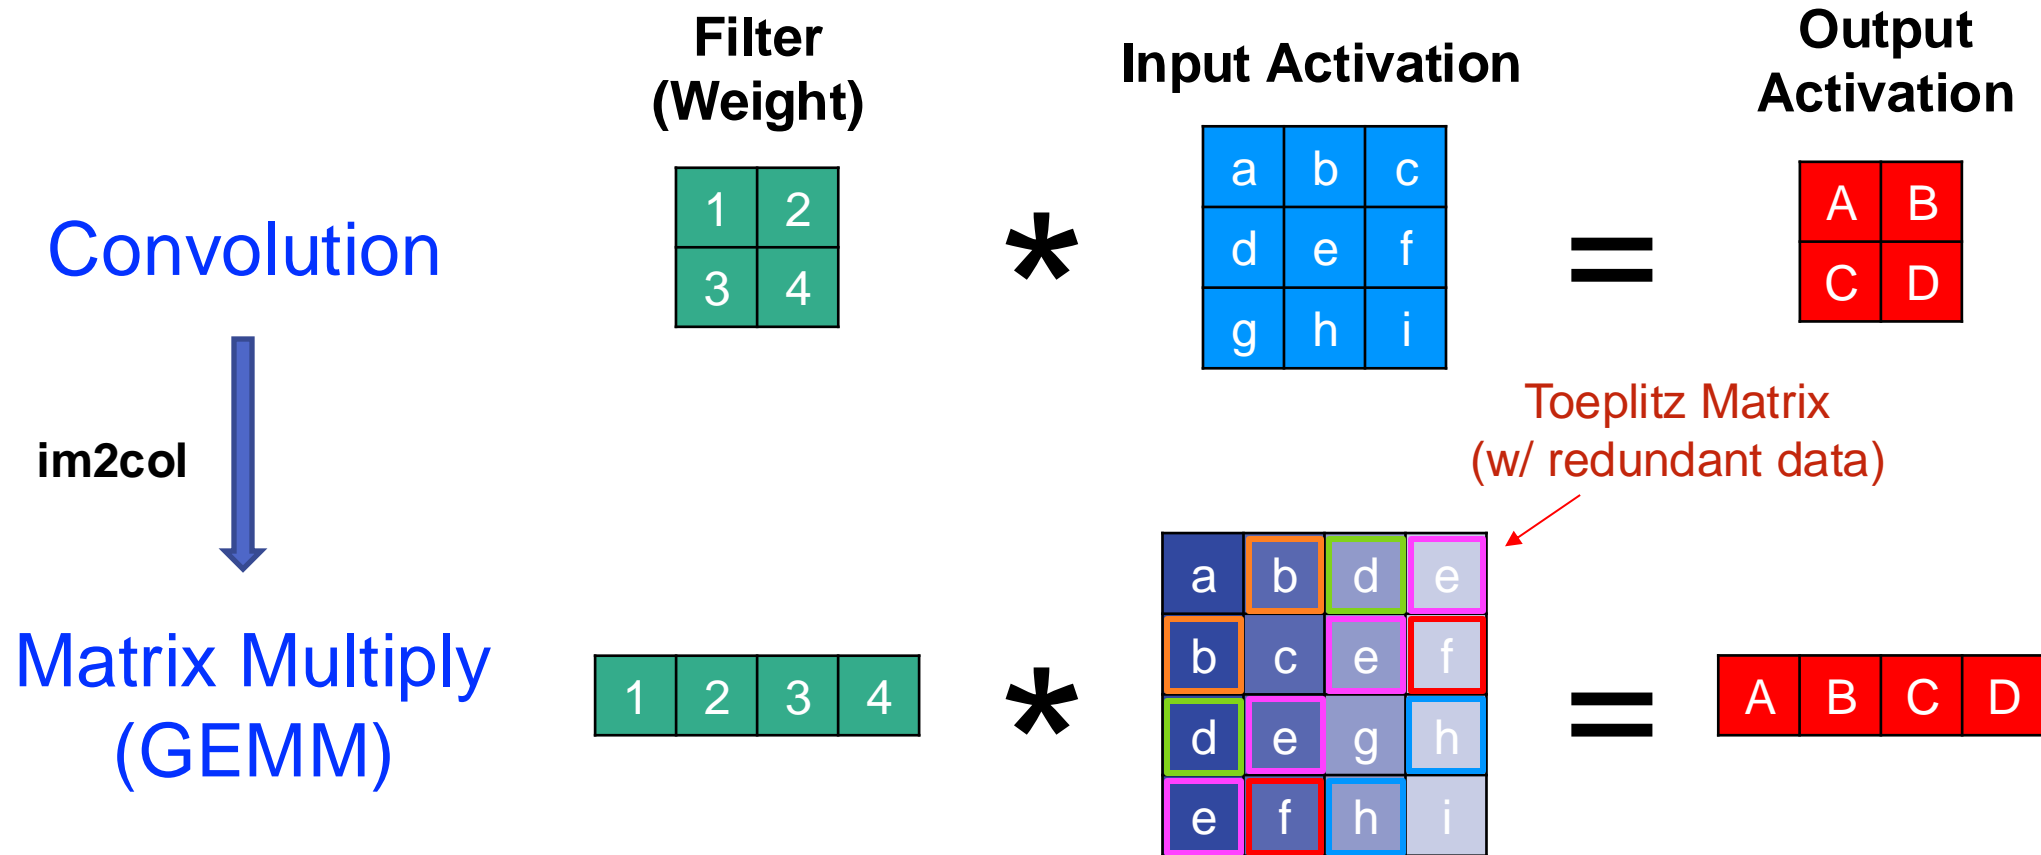


Batch Size = N



Recap: Mapping Convolution to Matrix Multiplication (Redundant Input Activations)

⦿ Converting convolution to GEMM via **im2col**





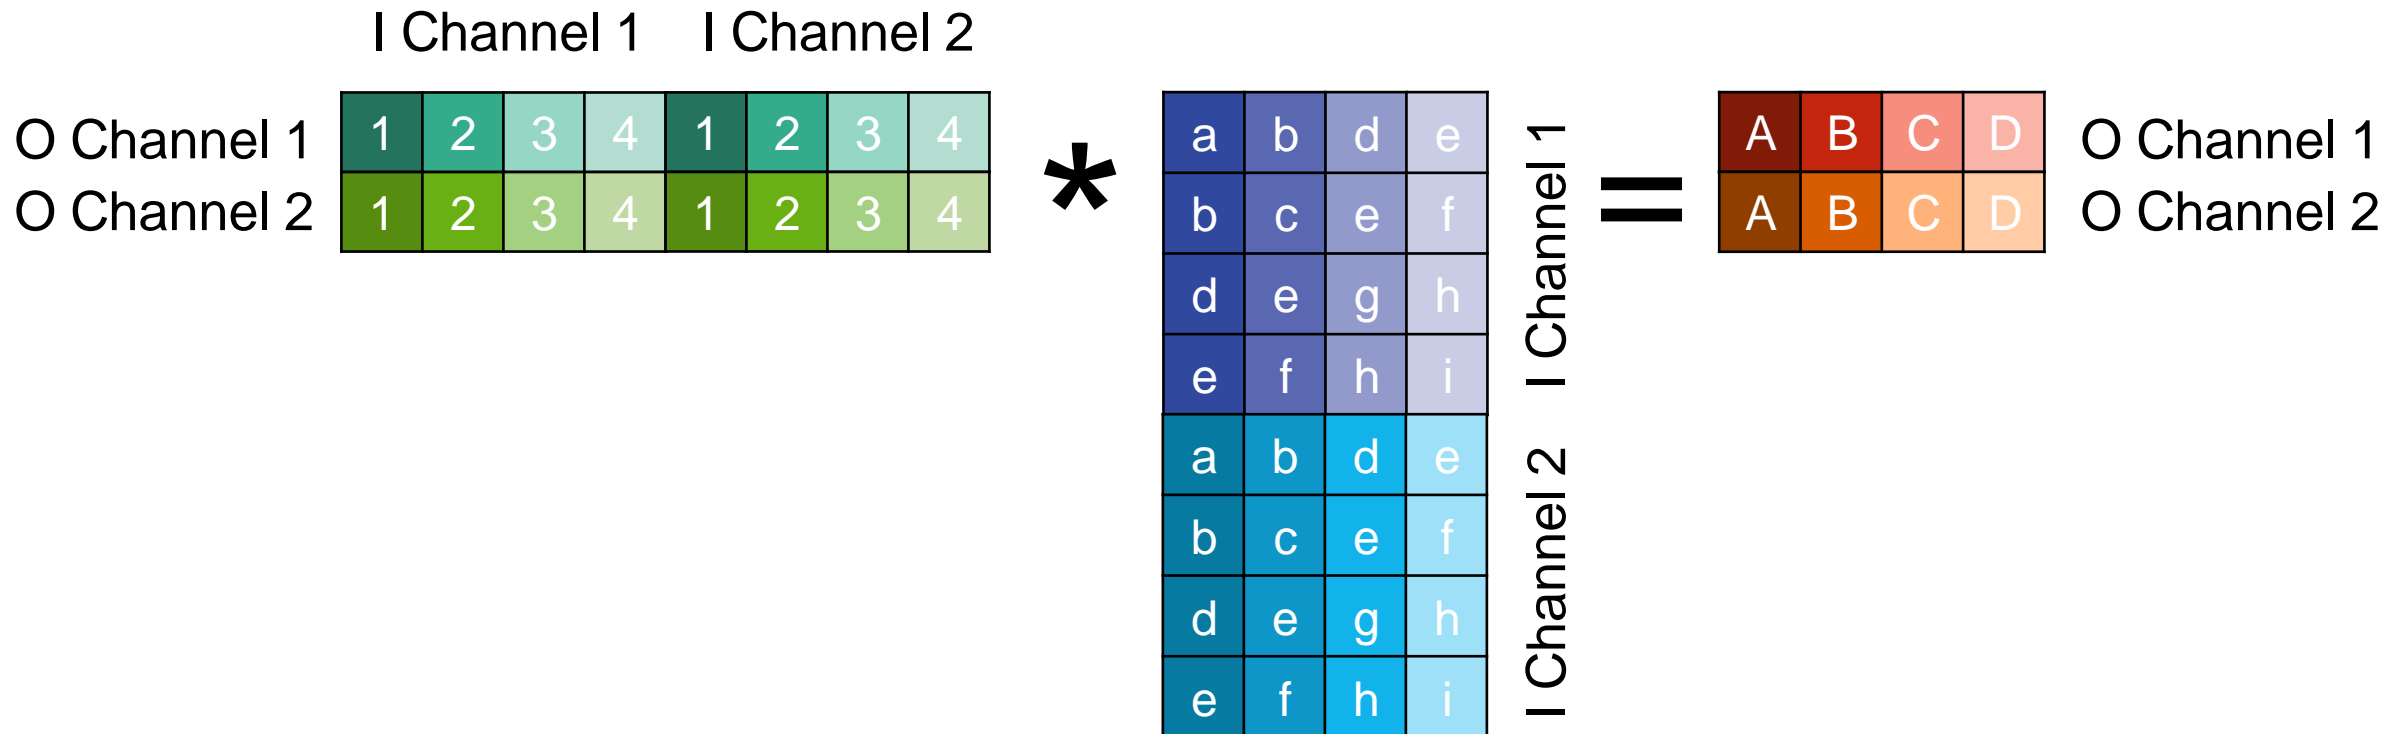
Mapping Convolution to Matrix Multiplication: Multiple Input/Output Channels

**Filter
(Weight)**

Input Activation


Toeplitz Matrix
(w/ redundant data)

**Output
Activation**





Mapping Convolution Layer to Matrix Multiplication



$$P = \frac{(H - R + U)}{U}, Q = \frac{(W - S + U)}{U}$$



Mapping Convolution to Matrix Multiplication (Redundant Weights)

Input Activation

a	b	c	d	e	f	g	h	i
---	---	---	---	---	---	---	---	---

*

Filter

1			
2	1		
	2		
3		1	
4	3	2	1
	4		2
		3	
		4	3
			4

=

Output Activation

A	B	C	D
---	---	---	---

--

 = Zero Value

Sparse Matrix






Tiling




Recap:

Naïve Matrix Multiplication for Fully-Connected Layer



Concept of Tiling: Tile-based Computation





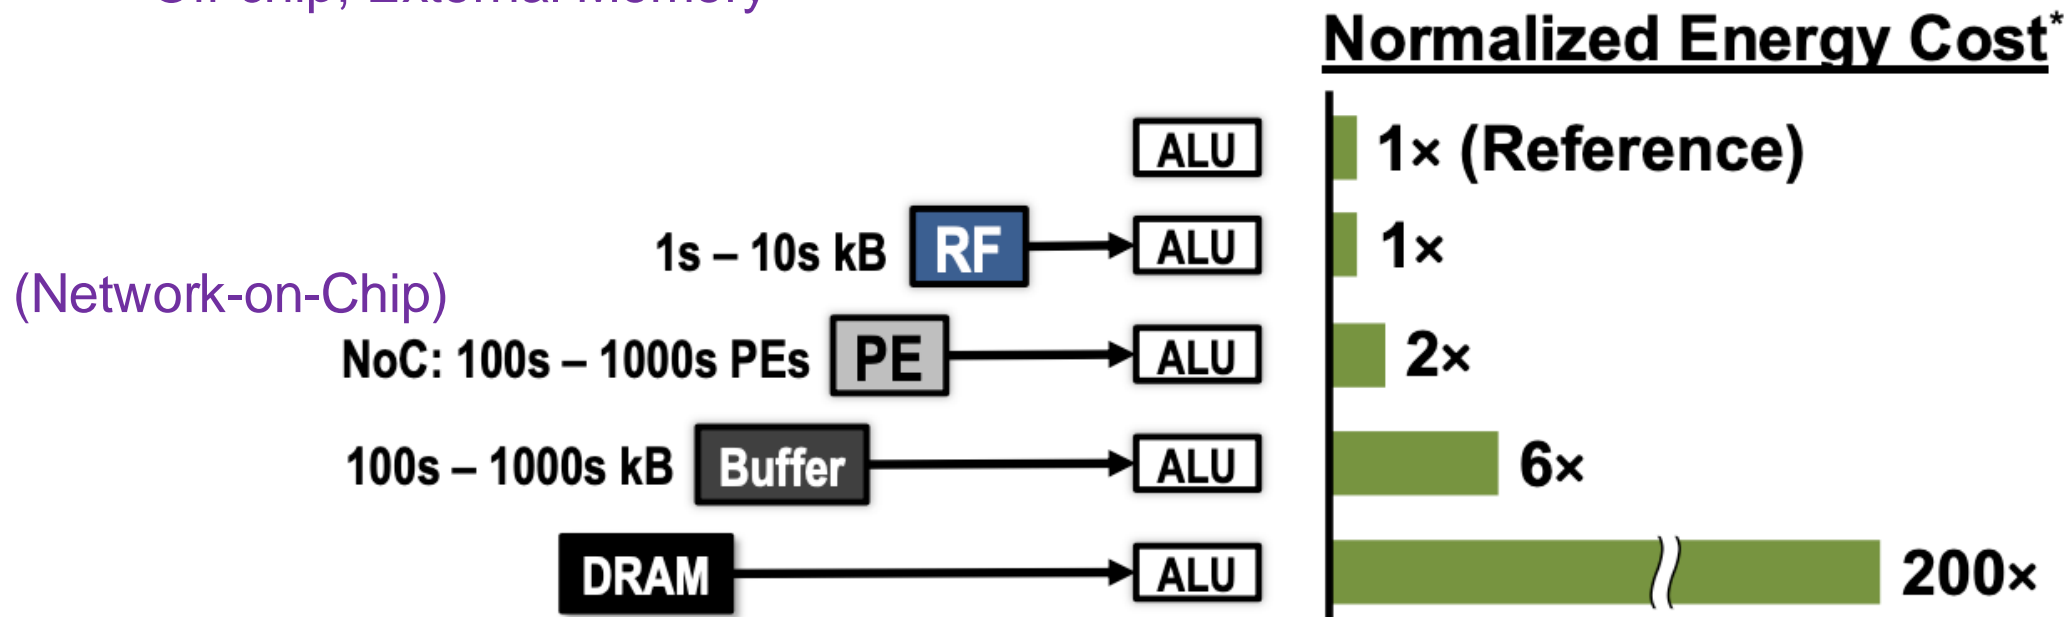
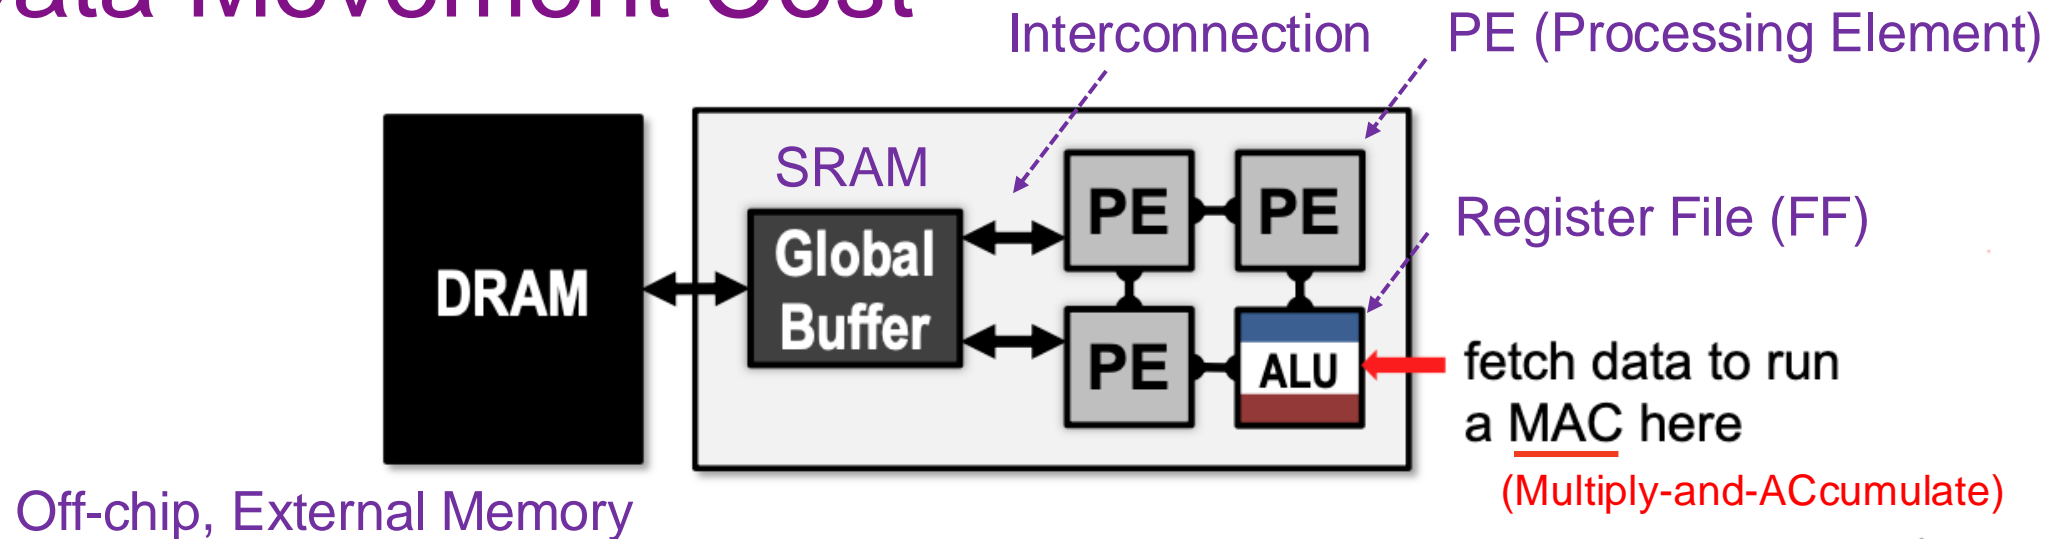
Tiling for Matrix Multiplication in Hardware

- ◎ Partition the matrix operation into small submatrices (tiles) to optimize hardware execution
 - ◆ Improves data reuse (locality) across different levels of memory hierarchy
 - ◆ Reduces memory bandwidth by keeping data closer to compute units
 - ◆ Allows for parallel execution on hardware accelerators
 - ◆ May generate subsequent partial results (**psum**) to be added
- ◎ Memory management
 - ◆ Not primarily managed by hardware **caches** (i.e., implicitly data orchestration)
 - ◆ Instead, explicitly managed local buffers (**scratchpads**) store and reuse tiles controlled explicitly by programmers or designers
 - ◆ Optimized data flow reduces costly global memory access



Data Locality and Reuse

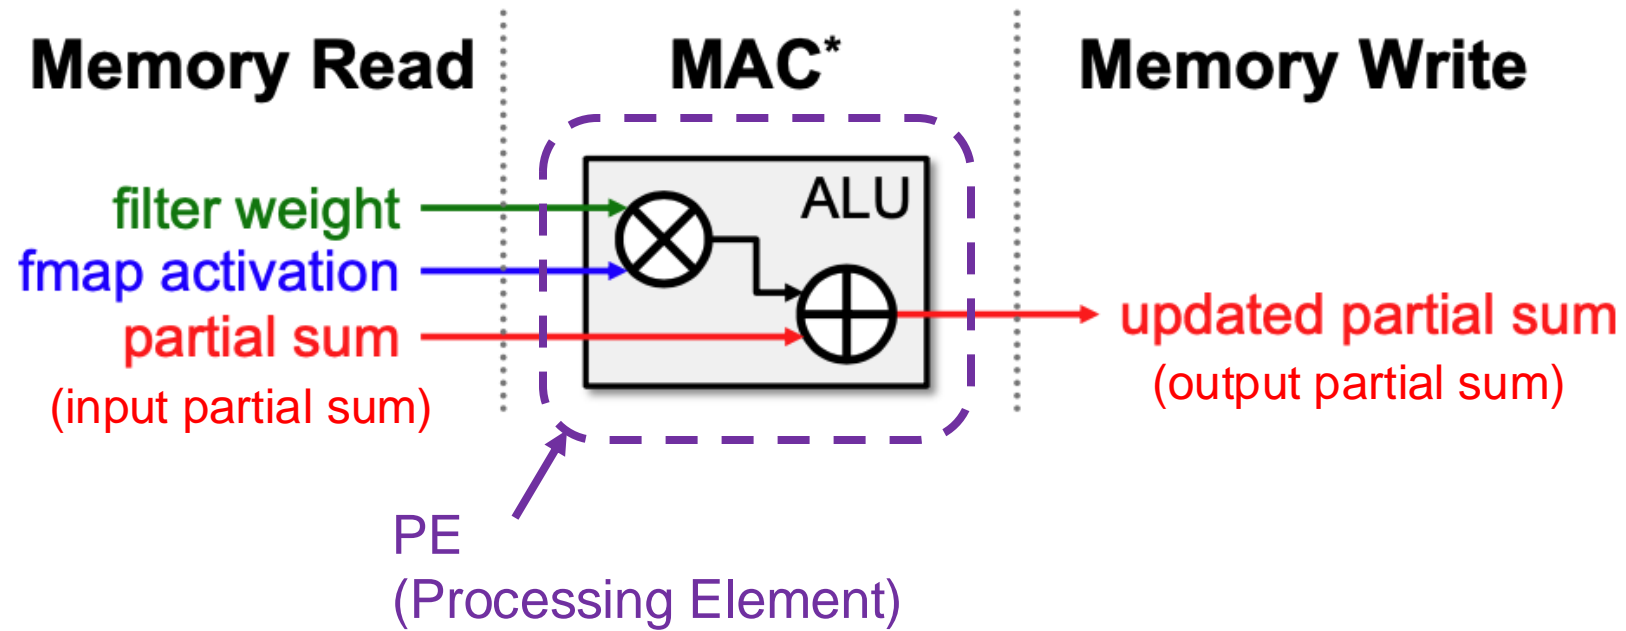
Data Movement Cost






Single MAC Unit for DNN

* MAC: Multiply-and-ACcumulate




Memory Access is the Bottleneck



* MAC: Multiply-and-ACcumulate

- ⦿ Worst case: all memory R/W are DRAM accesses
- ⦿ E.g., AlexNet has 724M MACs
 - 2896M DRAM accesses required
 - (extremely energy inefficient!)

Locality: Leverage Local Memory for Data Reuse



Extra levels of local memory hierarchy
Smaller, but Faster and more Energy-Efficient

(Usually more expensive with more levels)

- ⦿ **Temporal** reuse: the same data is used **more than once over time** by the same consumer




Temporal and Spatial Reuse

Temporal reuse

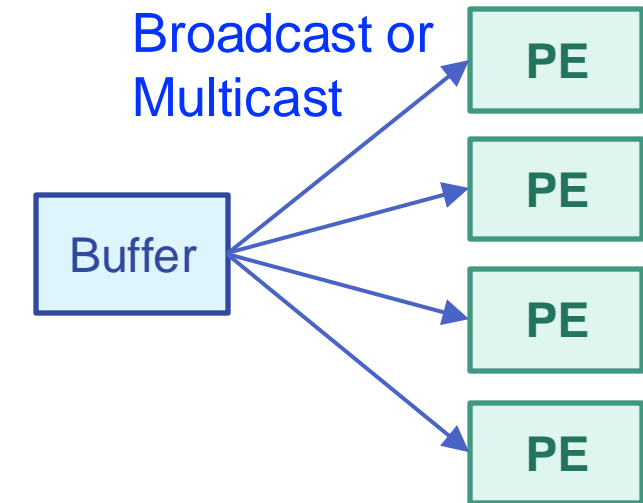
- ◆ The same data is used **more than once over time** by the same consumer

Memory Subsystem (Hierarchy)



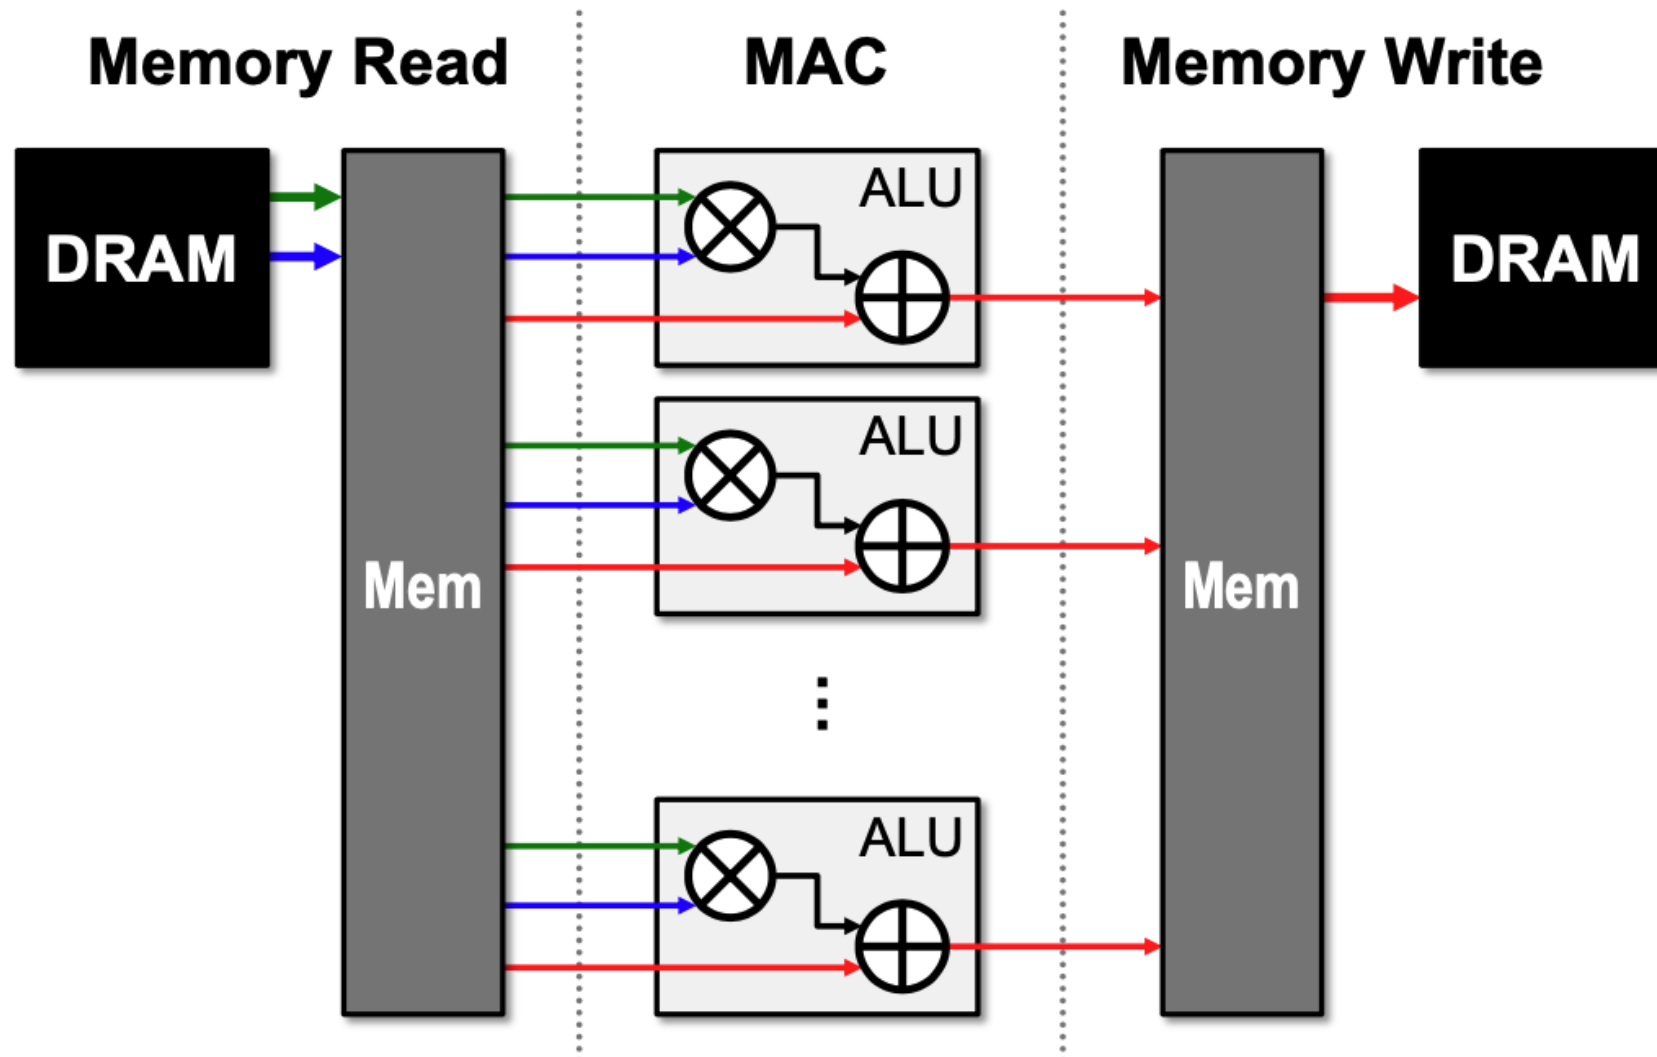
Spatial reuse

- ◆ The same data is used by **more than one consumer at different spatial locations** of the hardware





Leverage Parallelism for Higher Performance

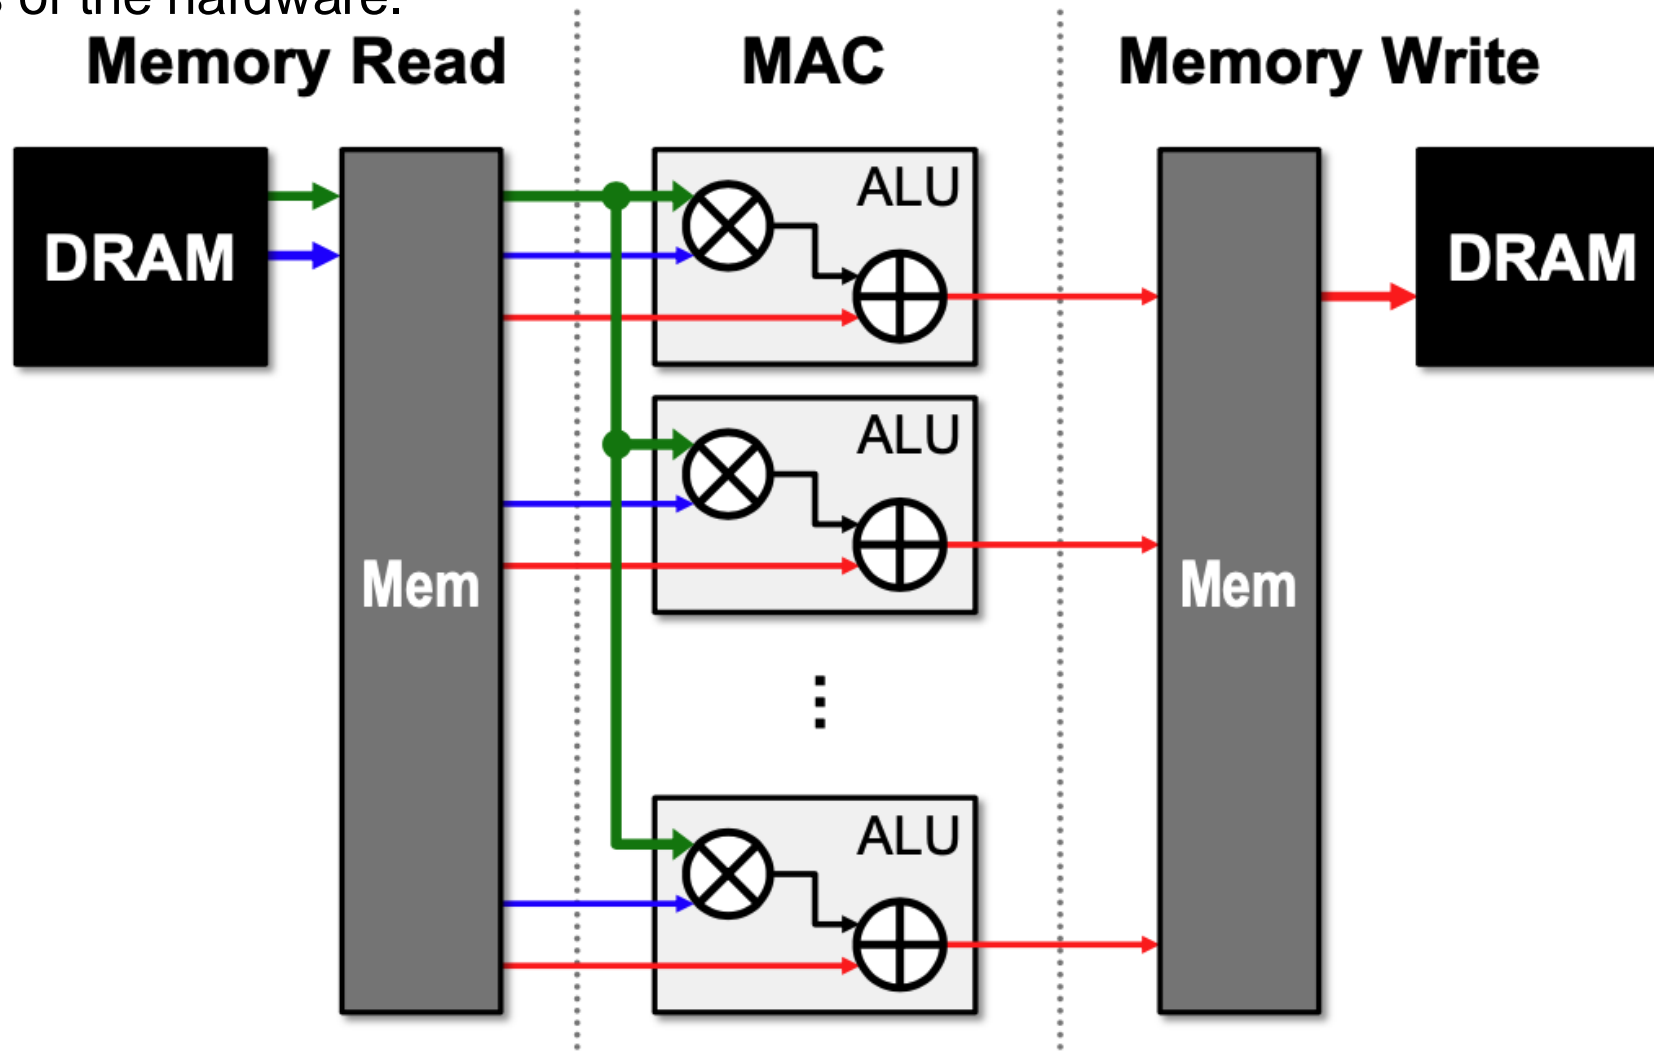


[Source: MIT Eyeriss Tutorial]



Leverage Parallelism with Spatial Data Reuse

- **Spatial reuse:** the same data is used by more than one consumer at different spatial locations of the hardware.



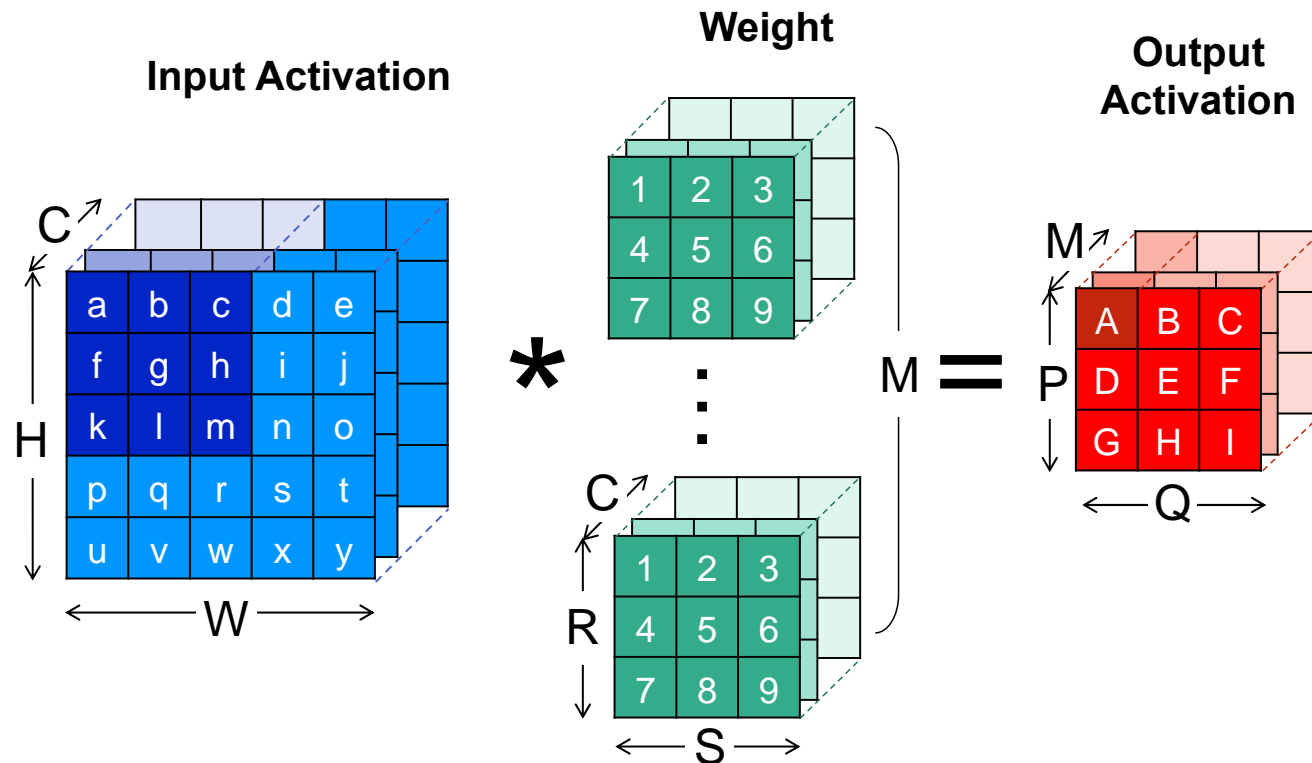
[Source: MIT Eyeriss Tutorial]



Data Reuse in DNN

Data Reuse in DNN

Batch Size = 1 ($N = 1$)



⊙ # MACs:

◆ $RSCPQM$

⊙ # Input Activations

◆ Size: HWC

◆ Max reuse: $\sim RSM$

⊙ # Weights

◆ Size: $RSCM$

◆ Max reuse: PQ

⊙ # Output Activations

◆ Size: PQM

◆ Max reuse: RSC

Dataflow and Mapping to Hardware


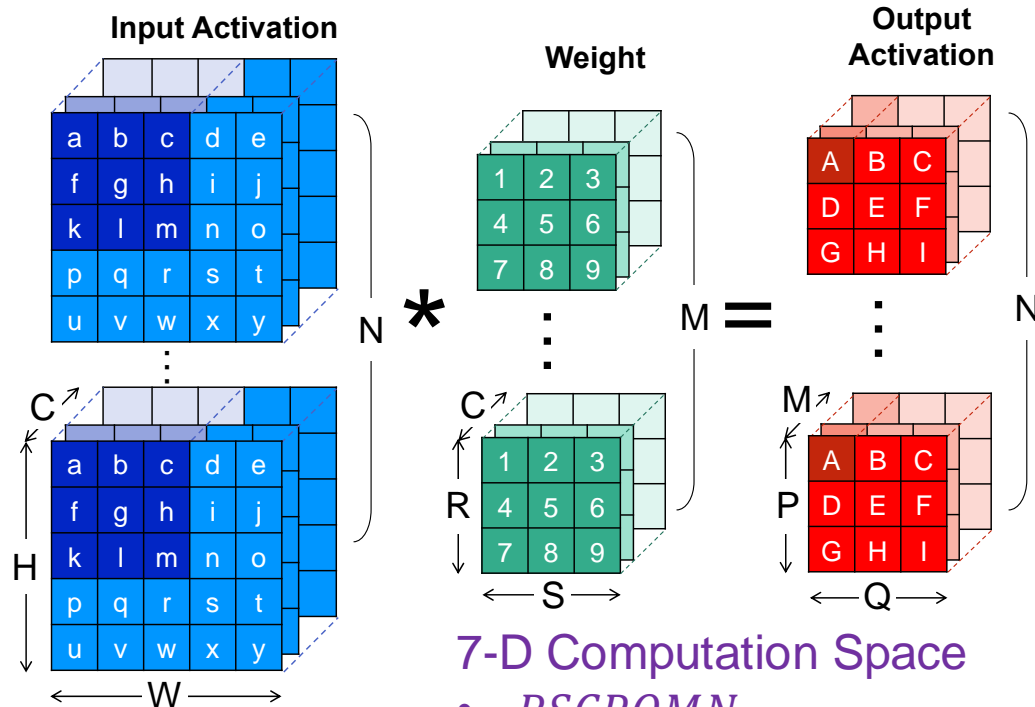
Millions of non-trivial mappings

7-dimensional Network Layer


Algorithmic Reuse

Hardware Reuse

1-d or 2-d (or 3-d or ...) Hardware



Data Movement via Interconnection
(Bus/crossbar/forwarding/NoC)



Temporal Data



Some Statistics for Your Reference

◎ AlexNet conv2

- ◆ Ifmap: 46K; Filter: 300K; Ofmap: 180K (total: 528K data)
- ◆ MAC: 224M

◎ MobileNet V1 conv2_1/dw

- ◆ Ifmap: 420K; Filter 9.2K, Ofmap: 400K (total: 829K data)
- ◆ MAC: 116M

◎ ResNet50 res5a_branch1

- ◆ Ifmap: 200K; Filter 2.1M, Ofmap: 100K (total: 2.4M data)
- ◆ MAC: 103M




Types of Data Reuse in DNN

For ideal data reuse, DRAM accesses in AlexNet can be reduced from 2896M to 61M

Convolutional Reuse


CONV layers only
(sliding window)



Reuse: **Activations**
Filter weights

Fmap Reuse


CONV and FC layers



Reuse: **Activations**

Filter Reuse

CONV and FC layers
(batch size > 1)



Reuse: **Filter weights**



Temporal Reuse




1-D Convolution with Temporal Reuse

$$\begin{bmatrix} 0 & 1 & 2 & 3 & 4 & 5 & 6 \end{bmatrix} * \begin{bmatrix} 0 & 1 & 2 & 3 \end{bmatrix} = \begin{bmatrix} 0 & 1 & 2 & 3 \end{bmatrix}$$


Example 1D convolution

Scheduling 1






Operation ordering 1: weight reuse distance = 4

Scheduling 2



Operation ordering 2: weight reuse distance = 1



* Only storage for weights is shown

Memory Hierarchy



Temporal Reuse to Improve Efficiency

- ⊙ Temporal reuse occurs when the same data value is used more than once by the same consumer (e.g., a PE)
 - ◆ By adding an intermediate memory level to the memory hierarchy of the hardware
- ⊙ Benefits
 - ◆ Less energy by accessing data from smaller memory level
 - ◆ May be faster too



Reuse Distance for Temporal Reuse

- ◎ **Reuse distance** of temporal reuse
 - ◆ Number of data accesses required by the consumer in between the accesses to the same data value
 - ◆ A function of the ordering of operations
- ◎ Storage capacity of intermediate memory limits the maximum reuse distance to be exploited
 - ◆ However, larger storage capacity implies more energy
- ◎ Reducing the reuse distance of one data type (**ifmap**, **weight**, or **ofmap**) often comes at the cost of increasing the reuse distance of other types
- ◎ Changing the processing order of compute can alter the reuse distance




Spatial Reuse




1-D Convolution with Spatial Reuse


* Only storage for weights is shown




(a) Memory hierarchy



(b) Operation ordering 1



(c) Operation ordering 2



(d) Operation ordering 3



Spatial Reuse to Improve Efficiency

- ◎ Spatial reuse occurs when the same data value is used by more than one consumer (e.g., a group of PEs) at different spatial locations of the hardware
 - ◆ By reading the data once from the source memory level and multicasting it to all consumers
- ◎ Benefits
 - ◆ Reducing # accesses to the source memory level
 - ▣ Less energy cost
 - ◆ Reducing the bandwidth required from the source memory level
 - ▣ Helps to keep the PEs busy (higher utilization) → higher performance



Reuse Distance for Spatial Reuse

◎ Reuse distance of spatial reuse

- ◆ The maximum number of data accesses in between any pair of consumers that access the same data value
- ◆ A function of the ordering of operations

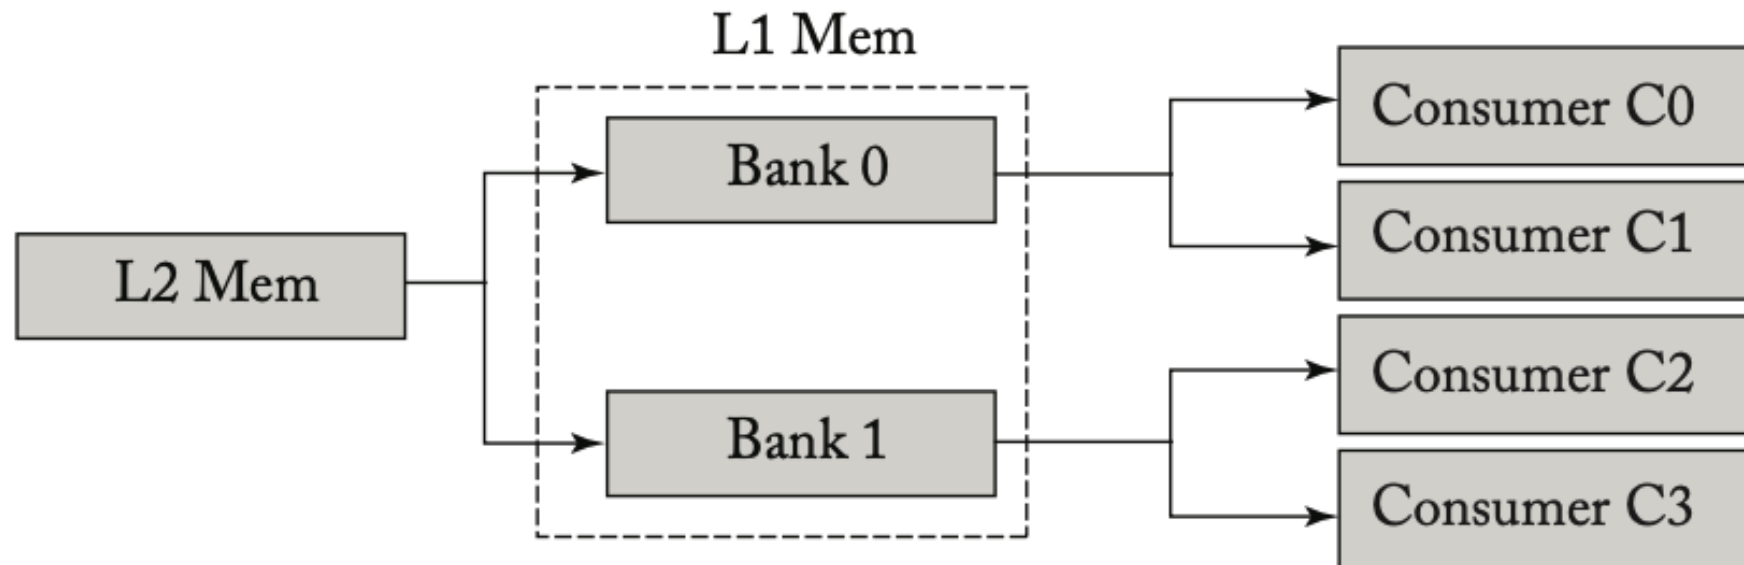
◎ If group of consumers have no storage capacity

- ◆ Spatial reuse by multicast in the same cycle
- ◆ Data need to be resent when it is needed in a different cycle



Physical Interconnection and Multi-Bank Memory


- Physical connectivity of the interconnection between L1 and consumers may limit the multicast from any banks in L1 to all consumers
- Data should be **multicast** from L2 to L1 → data duplication






Reducing Reuse Distance

Concept of Temporal Tiling to Reduce Reuse Distance of Partial Sums




Untiled Temporal Reuse



Tiled Temporal Reuse

3 slots of
intermediate
memory



* Only storage for weights is shown
Memory Hierarchy




Tiling

- ◎ The goal is to tile (partition) the data so that the reuse distance becomes smaller
 - ◆ However, it is not feasible to minimize the reuse distance for all data types simultaneously
- ◎ **Temporal tiling (tiling for temporal reuse)**
 - ◆ Reducing the reuse distance of specific data types to make it smaller than the storage capacity of a certain memory level in the memory hierarchy




Spatial Tiling (Tiling for Spatial Reuse)


- Reusing the same data value by as many consumers as possible
- Reducing the reuse distance so that one multicast can serve the maximum possible consumers given a fixed amount of storage capacity at each consumer



(a) No spatial reuse



(b) Medium degree of spatial reuse



(c) High degree of spatial reuse



Dataflow



Dataflow


◎ **Dataflow** refers to how data is processed within the hardware architecture

- ◆ Determines the path that data moves and how it is transformed and manipulated through the system
- ◆ Defines the execution order of the DNN operations in hardware
 - Computation order
 - Data movement order

◎ **Loop nest** is a compact way to describe the dataflow supported in hardware

- ◆ **for** (temporal for): describes the temporal execution order
- ◆ **parallel_for** (spatial for): describes the parallel execution

Weight-Stationary (WS) Dataflow and Output-Stationary (OS) Dataflow



```
for s in range(S):  
    for q in range(Q):  
        w = q+s;  
        OA[q] += IA[w] * W[s];
```


Weight-Stationary (WS) Dataflow

```
for q in range(Q):  
    for s in range(S):  
        w = q+s;  
        OA[q] += IA[w] * W[s];
```

Output-Stationary (OS) Dataflow



More Dataflow: Input-Stationary (IS)

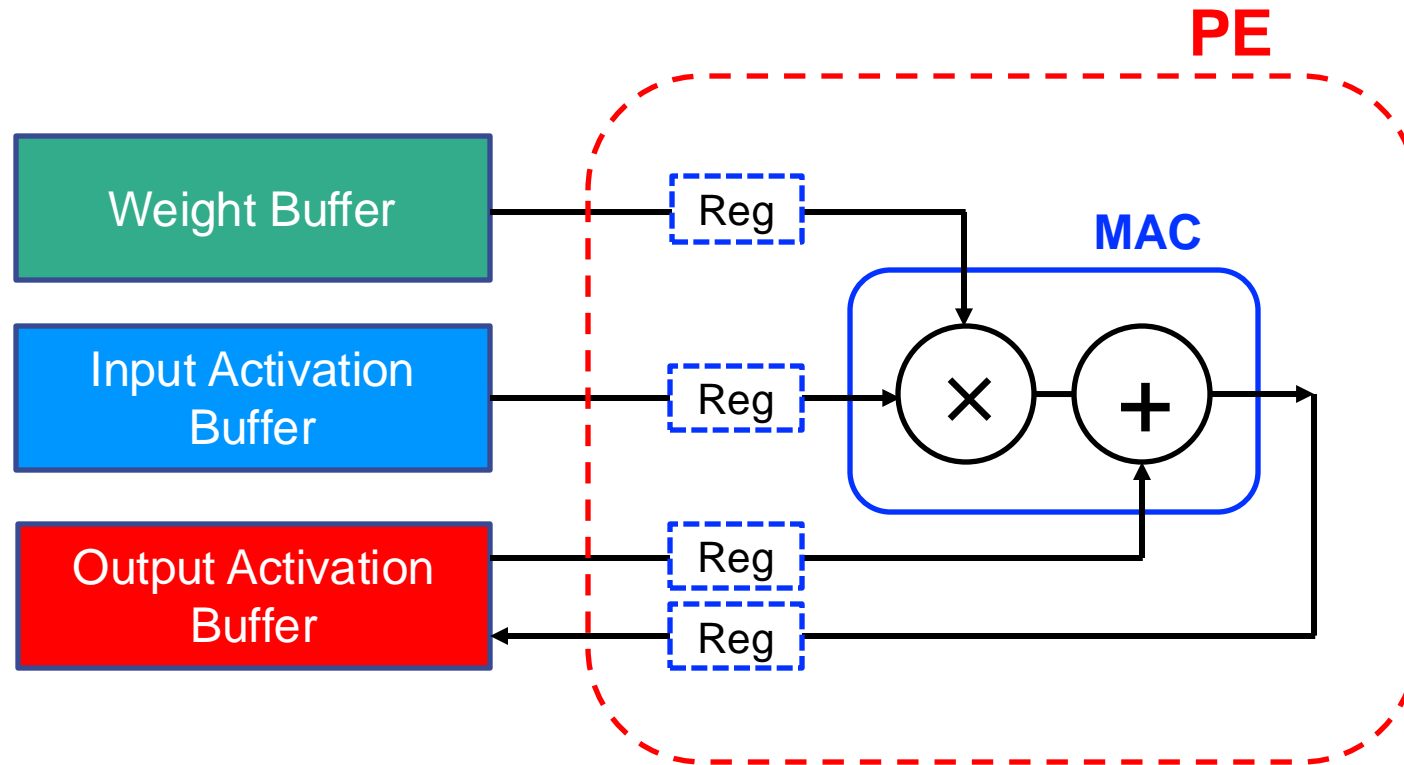


```
for w in range(W):  
    for s in range(S):  
        q = w-s;  
        OA[q] += IA[w] * W[s];
```

Input-Stationary (IS) Dataflow






Model of Single Processing Element (PE)



Space-Time Diagram for Weight-Stationary (WS) Dataflow

```
for s in range(S):  
    for q in range(Q):  
        w = q+s  
        OA[q] += IA[w] * W[s];
```


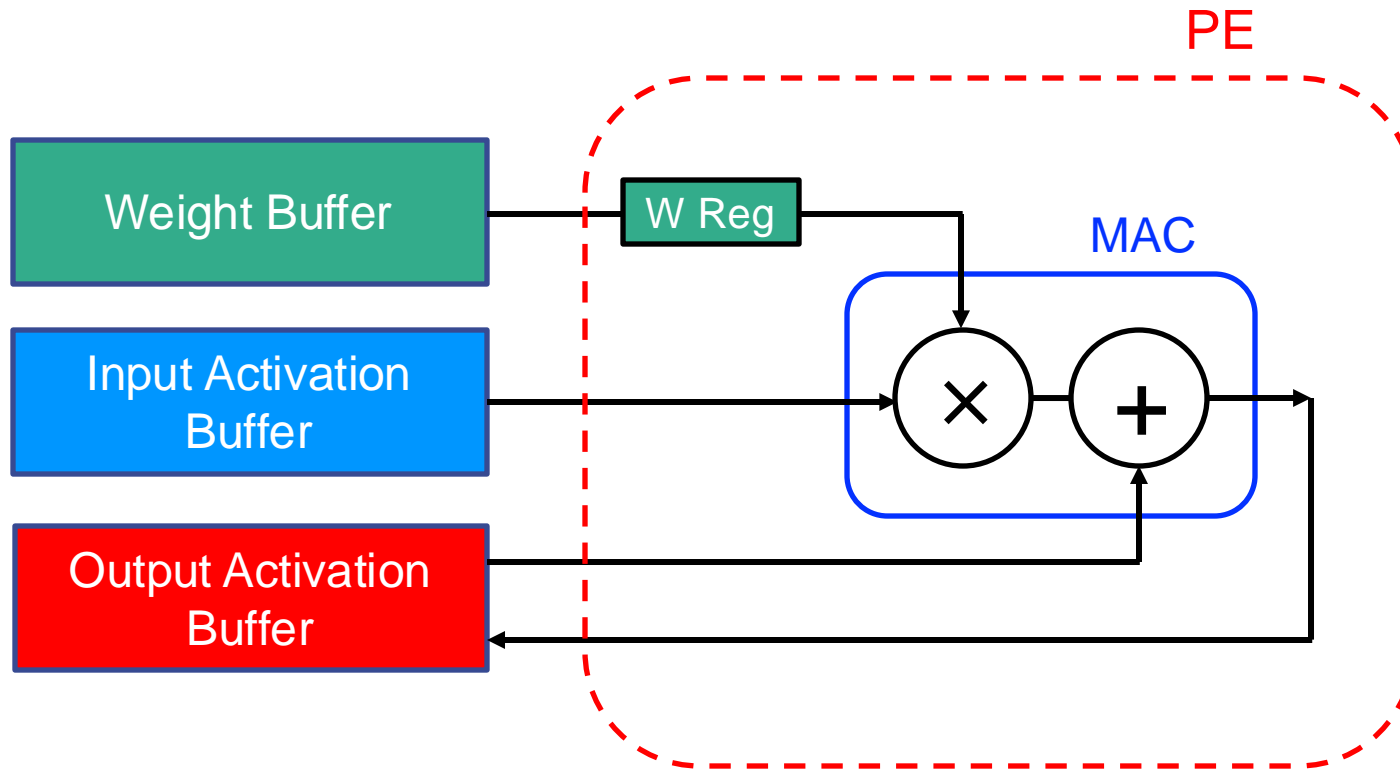
$$Q = 9, W = 12, S = 4$$






Space-Time Diagram for Weight-Stationary (WS) Dataflow


Observation: Single **weight** reused Q times




Weight-Stationary Dataflow



Weight-Stationary Dataflow






Weight-Stationary Dataflow



Space-Time Diagram for Output-Stationary (OS) Dataflow

```
for q in range(Q):  
    for s in range(S):  
        w = q+s  
        OA[q] += IA[w] * W[s];
```


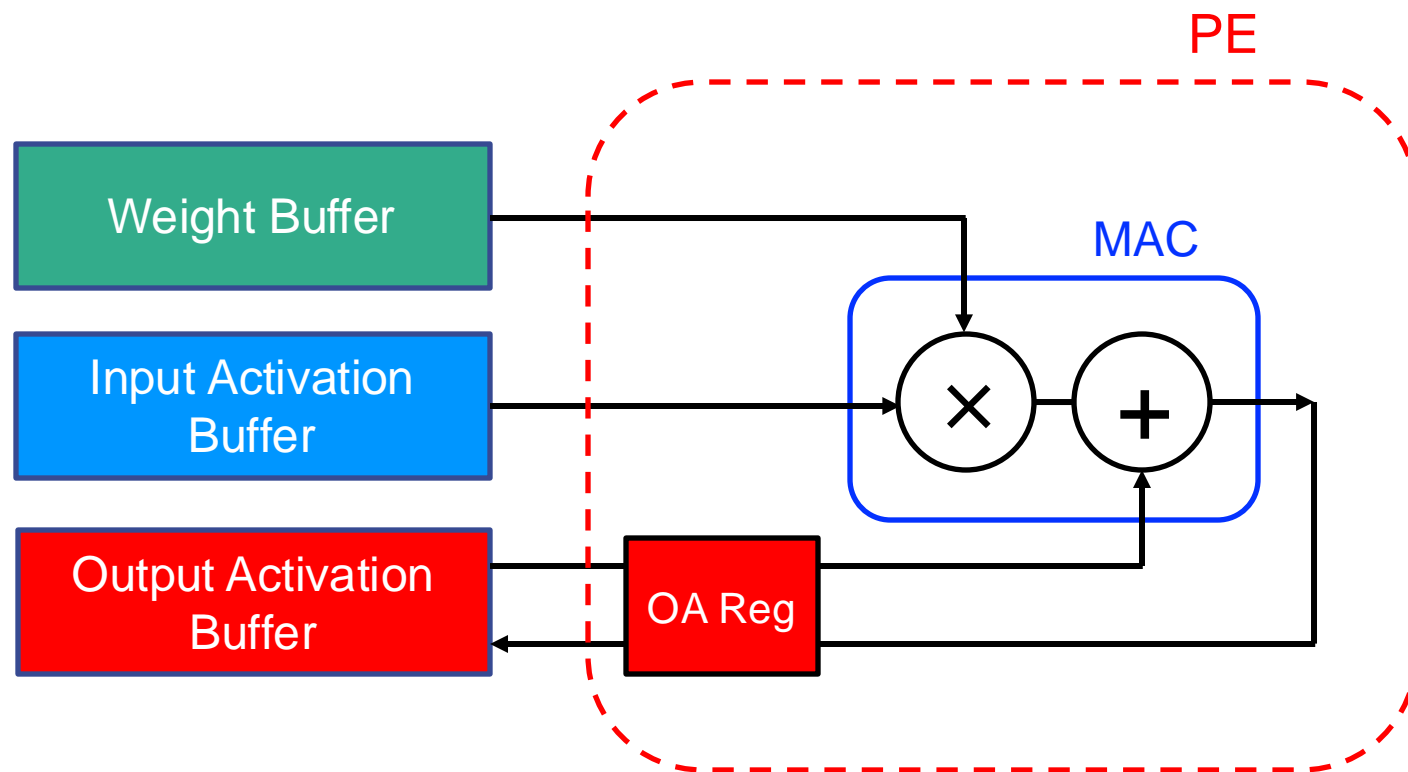
$$Q = 9, W = 12, S = 4$$




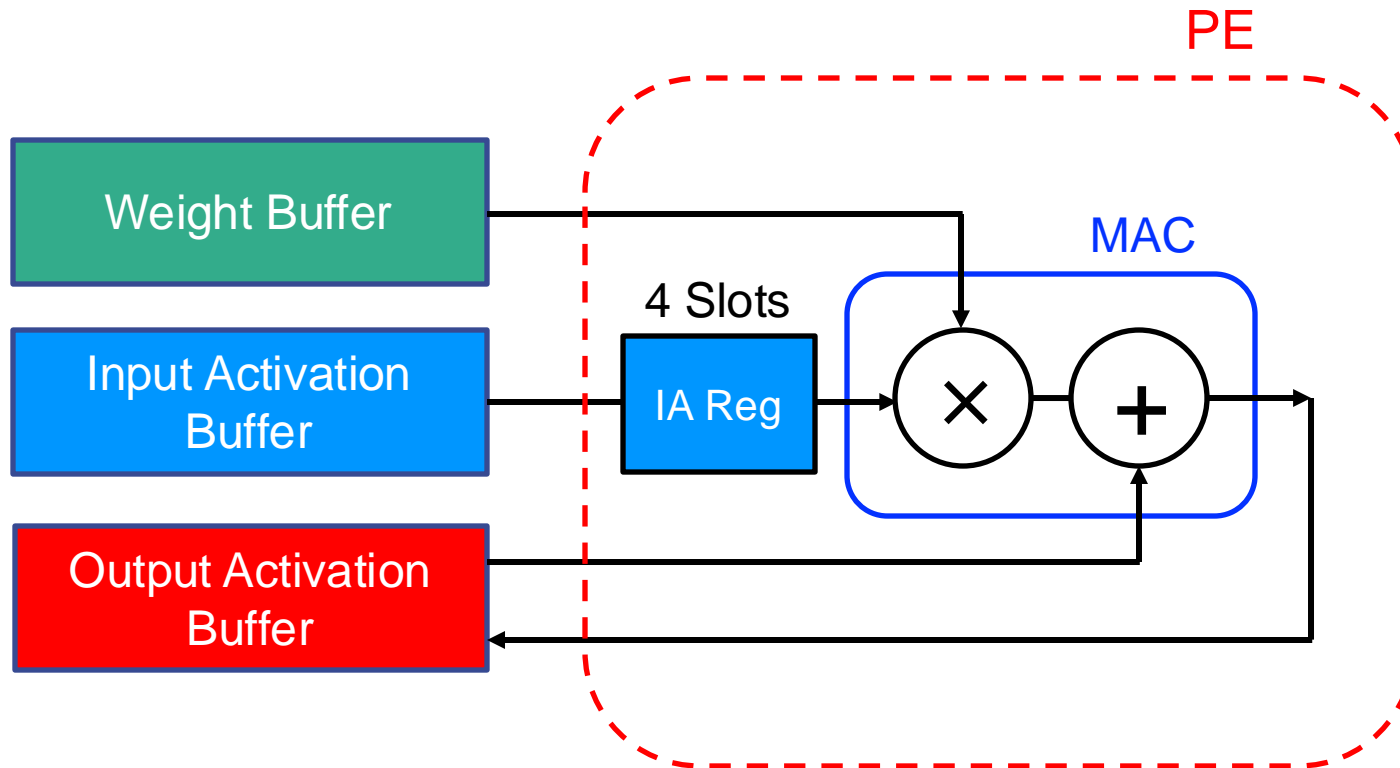


Space-Time Diagram for Output-Stationary (OS) Dataflow


Observation: Single **output** reused S times




Sliding Window in Output-Stationary




Output-Stationary Dataflow




Output-Stationary Dataflow



Output-Stationary Dataflow



Output-Stationary Dataflow






Tiled Loop Nest for Dataflow



Tiled Loop Nest for Weight-Stationary Dataflow






$Q = 9, W = 12, S = 4$; S is tiled into $S1 = 2, S0 = 2$

```
for s1 in range(S1):  
    for q in range(Q):  
        parallel_for s0 in range(S0):  
            s = s1*S0 + s0;  
            w = q+s;  
            OA[q] += IA[w] * W[s];
```

Space-Time Diagrams of Tiled WS Dataflow

```
for s1 in range(S1):  
    for q in range(Q):  
        parallel_for s0 in range(S0):  
            s = s1*S0 + s0; w = q+s;  
            OA[q] += IA[w] * W[s];
```

$Q = 9, W = 12, S = 4$; S is tiled into $S1 = 2, S0 = 2$





Summary

- ◎ Parallelizing the compute may reduce the latency
- ◎ Minimizing data movement is the key to high energy efficiency for DNN accelerators
 - ◆ Maximizing data reuse at the higher level of memory hierarchy
- ◎ Dataflow taxonomy:
 - ◆ Output Stationary: minimize movement of **psums**
 - ◆ Weight Stationary: minimize movement of **weights**
 - ◆ Input Stationary: minimize movement of **inputs**
- ◎ Loop nest provides a compact way to describe properties of a dataflow
 - ◆ E.g., **data tiling** in multi-level storage and temporal/spatial processing