



CS5120 VLSI System Design, Spring 2025

DNN Mapping Part 2: Tiling for Hardware Structure

黃稚存

Chih-Tsun Huang

cthuang@cs.nthu.edu.tw



國立清華大學
NATIONAL TSING HUA UNIVERSITY

資訊工程學系
Computer Science

Lecture 07



聲明

- ◎本課程之內容 (包括但不限於教材、影片、圖片、檔案資料等)，僅供修課學生個人合理使用，非經授課教師同意，不得以任何形式轉載、重製、散布、公開播送、出版或發行本影片內容 (例如將課程內容放置公開平台上，如 Facebook, Instagram, YouTube, Twitter, Google Drive, Dropbox 等等)。如有侵權行為，需自負法律責任。



Outline


- ① Hardware Resource Constraints
- ① Mapping Space
- ① Tuning of DNN Mapping



Recap: Matrix Multiplication for Fully-Connected Layer

⊙ Consider $A * B = C$


- ◆ Can be generalized to convolution





Recap: Loop Nest for Matrix Multiplication

⊙ Consider $A * B = C$

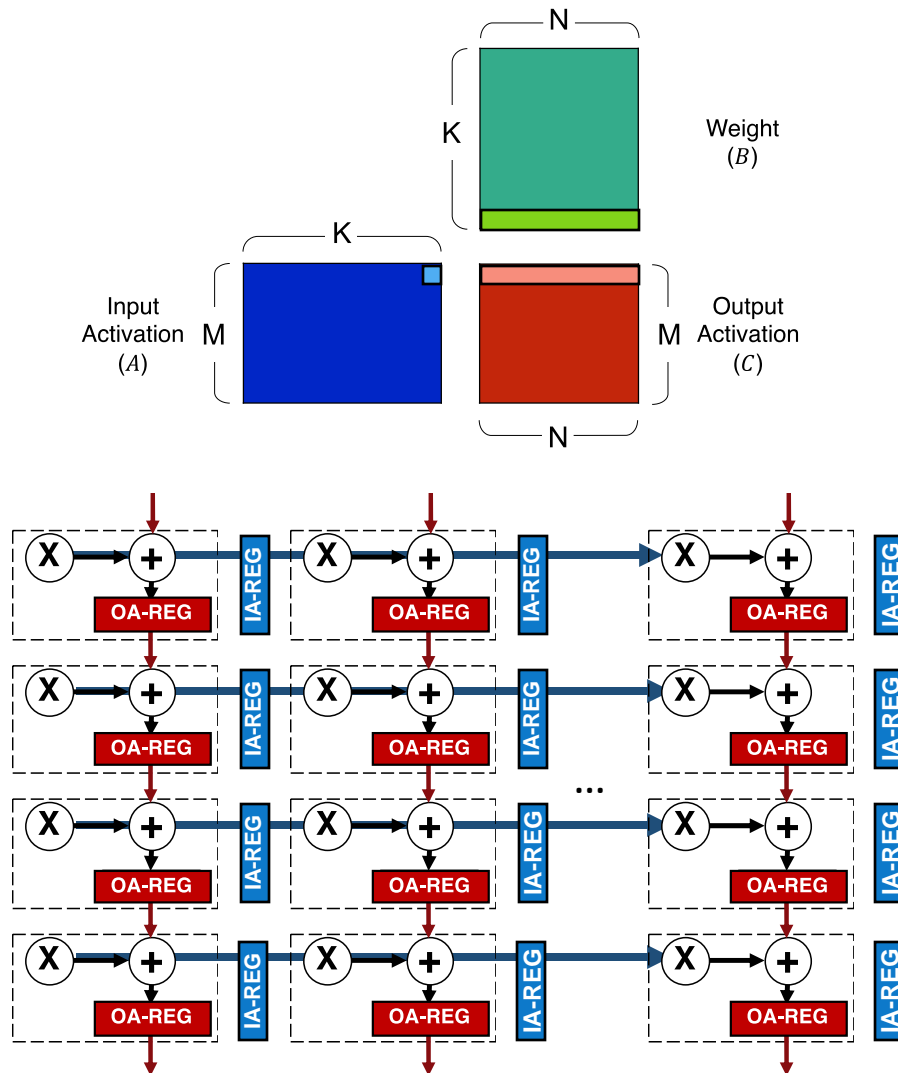


```
for (m=0; m<M; m++) {  
    for (n=0; n<N; n++) {  
        OA[n,m] = 0;  
        for (k=0; k<K; k++) {  
            OA[n,m] += IA[m, k] * W[k, n];  
        }  
        OA[n,m] = Activation(OA[n,m]);  
    }  
}
```

For each output activation

Reduction

Recap: Mixed Datapath Optimization: TPU



```
for (m=0; m<M; m++) {  
  parallel_for (n=0; n<N; n++) {  
    OA[n,m] = 0;  
    parallel_for (k=0; k<K; k++) {  
      OA[n,m] += IA[m,k] * W[k,n];  
    }  
    OA[n,m] = Activation(OA[n,m]);  
  }  
}
```

● Systolic accumulation

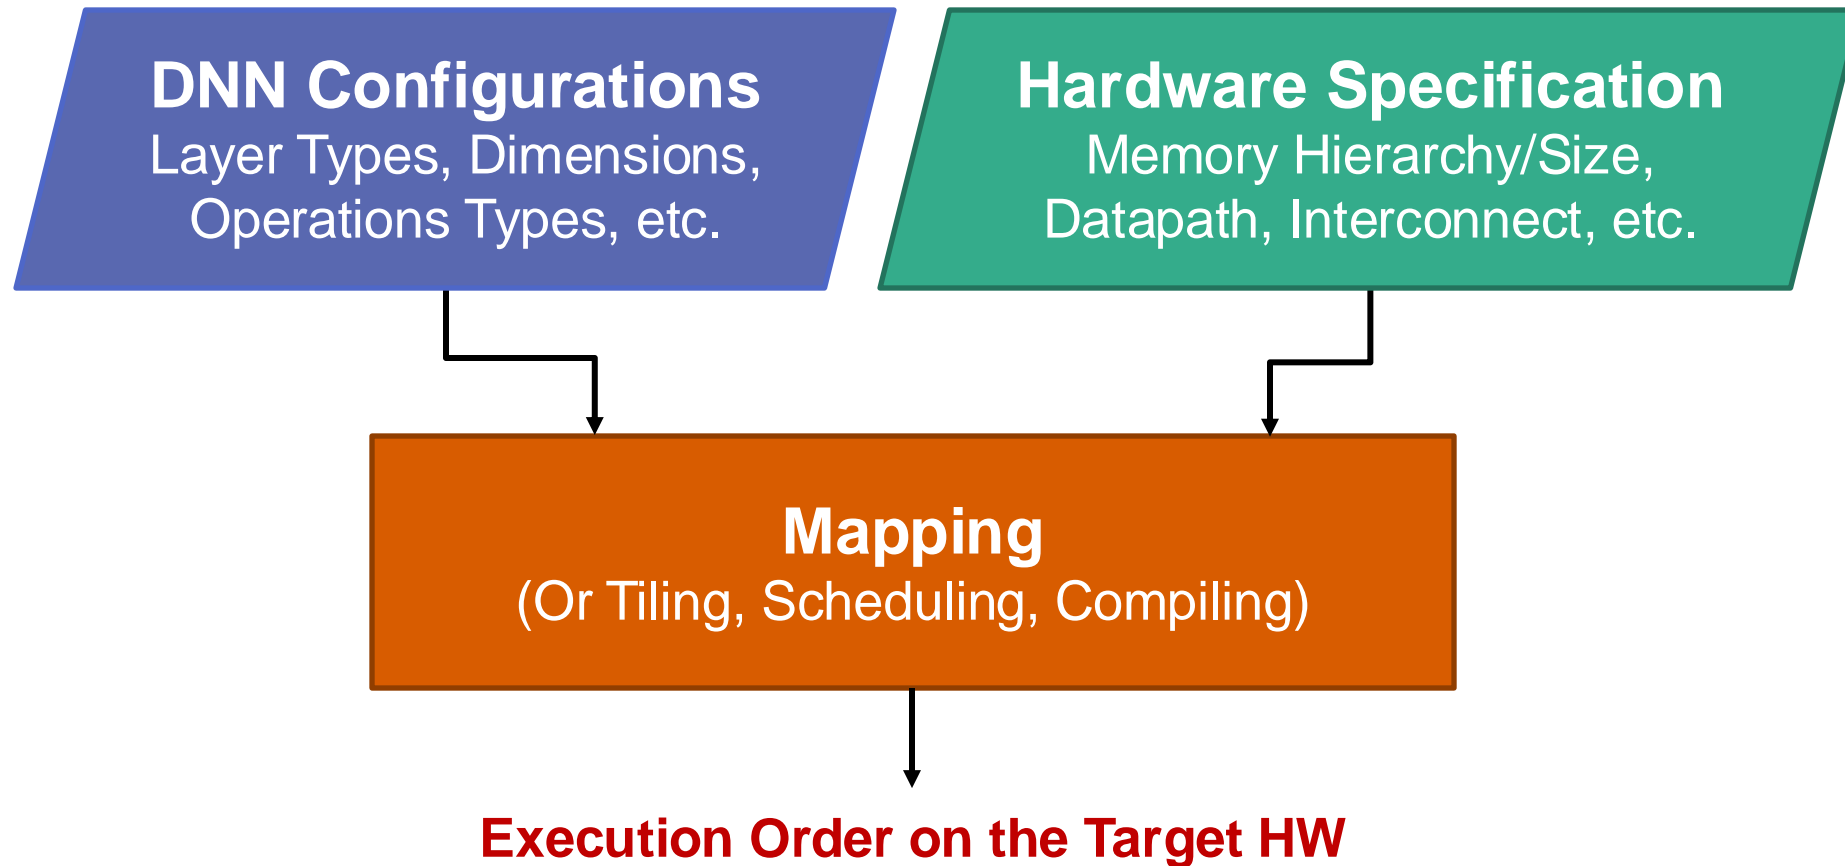
● Systolic multicast

→ Area vs. scalability

→ Latency vs. pipeline throughput



DNN Mapping Problem




- ⦿ Mapping objective: efficient for the performance (latency) and/or energy, etc.




Hardware Resource Constraints

Constraint 1: Systolic Array Size < K and/or N

⦿ DNN Configuration: Matrix Multiplication



⦿ Hardware Specification: Systolic array size: 2x2



⦿ Notation: $K?/N?$ \rightarrow loop bounds

```

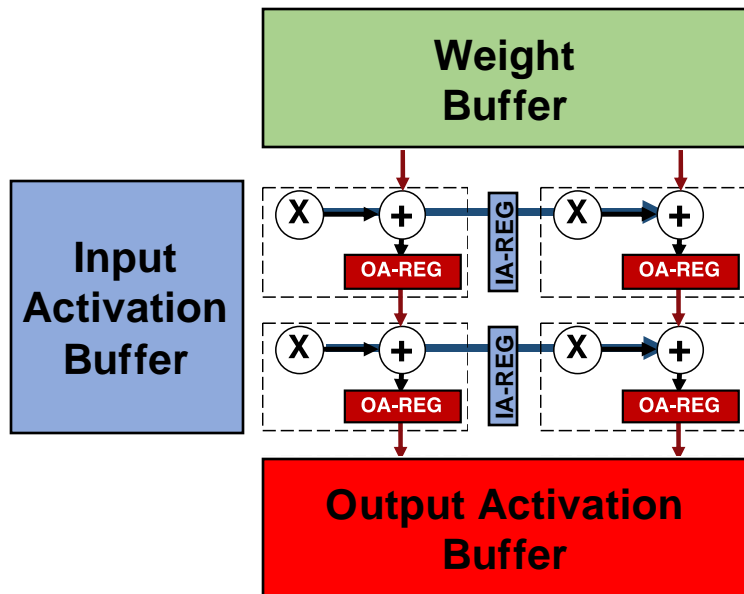

for (m=0; m<M; m++) {
    for (n1=0; n1<N1; n1++) {
        OA[n1*N0:(n1+1)*N0,m] = 0;
        for (k1=0; k1<K1; k1++) {
            parallel_for (n0=0; n0<N0; n0++) {
                parallel_for (k0=0; k0<K0; k0++) {
                    OA[n1*N0+n0,m] += IA[m,k1*K0+k0]
                        * W[k1*K0+k0,n1*N0+n0];
                }
            }
        }
    }
}
    
```

Temporal Tiling (indicated by a purple bracket on the right side of the code block)

Spatial Tiling (indicated by a purple bracket on the left side of the code block)

Constraint 2: Weight Buffer Size $< K * N$

(1/3)



Hardware Specification:

- ◆ Systolic array size: 2×2
- ◆ Explicit data movement (or data orchestration)

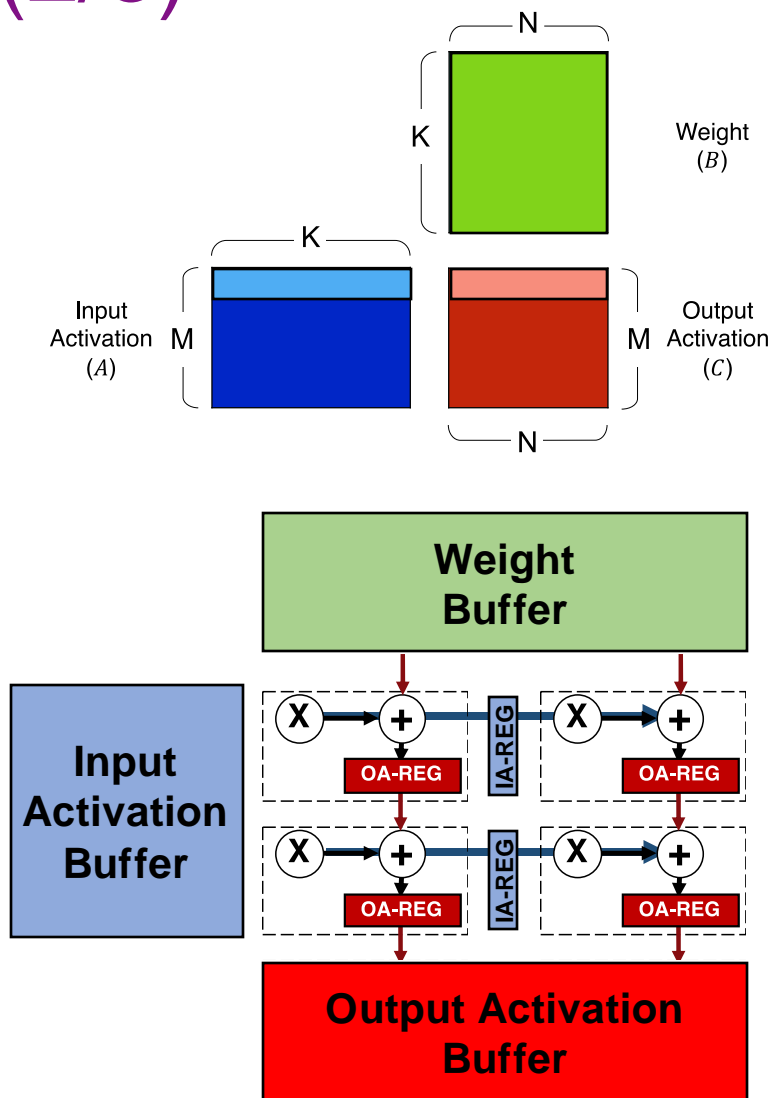
```
for (m=0; m<M; m++) {
```

```
    for (n1=0; n1<N1; n1++) {
        OA[n1*N0:(n1+1)*N0,m] = 0;
        for (k1=0; k1<K1; k1++) {
            parallel_for (n0=0; n0<N0; n0++) {
                parallel_for (k0=0; k0<K0; k0++) {
                    OA[n1*N0+n0,m] += IA[m,k1*K0+k0]
                                   * W[k1*K0+k0,n1*N0+n0];
                }
            }
        }
    }
```

```
}
```

Constraint 2: Weight Buffer Size $< K * N$

(2/3)



Hardware Specification:

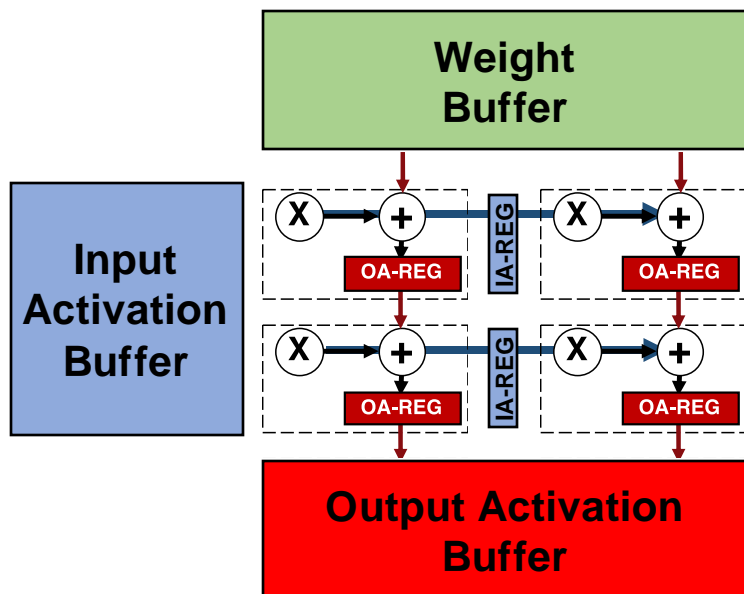

- ◆ Systolic array size: 2×2
- ◆ Explicit data movement (or data orchestration)
- ◆ Weight, input activation, output activation **buffer sizes**

```

for (m=0; m<M; m++) {
    mvin(W[0:K,0:N]);           // W buffer  $\geq K*N$ 
    mvin(IA[m:m+1,0:K]);        // IA buffer  $\geq 1*K$ 
    for (n1=0; n1<N1; n1++) {
        OA[n1*N0:(n1+1)*N0,m] = 0;
        for (k1=0; k1<K1; k1++) {
            parallel_for (n0=0; n0<N0; n0++) {
                parallel_for (k0=0; k0<K0; k0++) {
                    OA[n1*N0+n0,m] += IA[m,k1*K0+k0]
                                     * W[k1*K0+k0,n1*N0+n0];
                }
            }
        }
    }
    mvout(OA[0:N,m:m+1]);       // OA buffer  $\geq 1*N$ 
}
    
```

Constraint 2: Weight Buffer Size $< K * N$

(3/3)

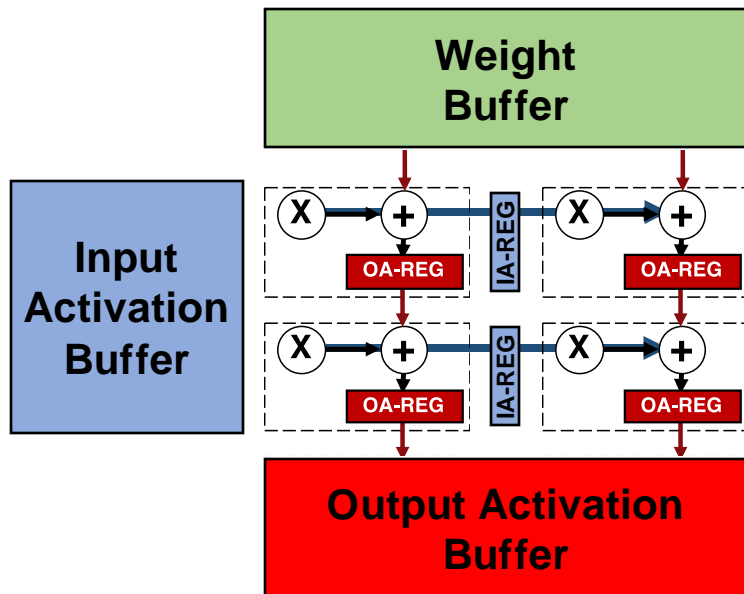



```

for (m=0; m<M; m++) {
    // IA buffer: 1*K
    mvin(IA[m:m+1,0:K]);
    for (n2=0; n2<N2; n2++) {
        // W buffer: N1*N0*K < K*N
        mvin(W[0:K,n2*N1*N0:(n2+1)*N1*N0]);
        OA[n2*N1*N0:(n2+1)*N1*N0,m:m+1]=0;
        for (n1=0; n1<N1; n1++) {
            for (k1=0; k1<K1; k1++) {
                parallel_for (n0=0; n0<N0; n0++) {
                    parallel_for (k0=0; k0<K0; k0++) {
                        OA[n2*N1*N0+n1*N0+n0,m]
                            += IA[m,k1*K1+k0]
                               * W[k1*K0+k0,n2*N1*N0+n1*N0+n0];
                    }
                }
            }
        }
        mvout(OA[n2*N1*N0:(n2+1)*N1*N0,m:m+1]);
    }
}

```

Constraint 3: Input Buffer Size > 1 * K

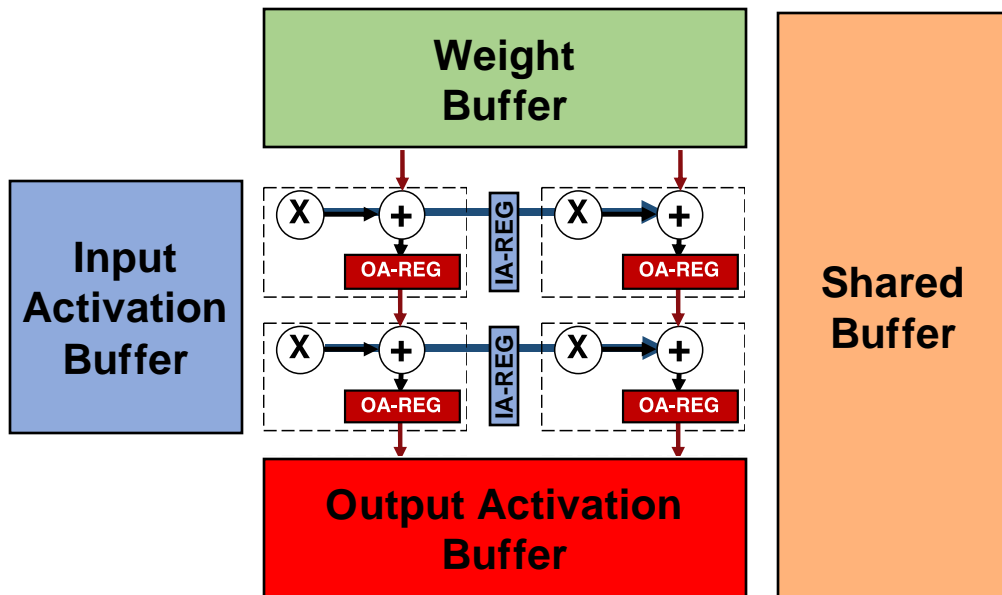



```

for (m2=0; m2<M2; m2++) {
    // IA buffer:  $M1 \times K > 1 \times K$ 
    mvin(IA[m2*M1:(m2+1)*M1,0:K]);
    for (n2=0; n2<N2; n2++) {
        // W buffer:  $N1 \times N0 \times K < K \times N$ 
        mvin(W[0:K,n2*N1*N0:(n2+1)*N1*N0]);
        OA[n2*N1*N0:(n2+1)*N1*N0,m2*M1:(m2+1)*M1]=0;
        for (m1=0; m1<M1; m1++) {
            for (n1=0; n1<N1; n1++) {
                for (k1=0; k1<K1; k1++) {
                    parallel_for (n0=0; n0<N0; n0++) {
                        parallel_for (k0=0; k0<K0; k0++) {
                            OA[n2*N1*N0+n1*N0+n0,m2*M1+m1]
                                += IA[m2*M1+m1,k1*K1+k0]
                                    * W[k1*K0+k0,n2*N1*N0+n1*N0+n0];
                        }
                    }
                }
            }
        }
        mvout(OA[n2*N1*N0:(n2+1)*N1*N0, m2*M1:(m2+1)*M1]);
    }
}

```

Constraint 4: Adding Another Shared Buffer



```
for (m3=0; m3<M3; m3++) {
  for (n3=0; n3<N3; n3++) {
    // Shared buffer blocking
```

```
for (m2=0; m2<M2; m2++) {
  // IA buffer stores: M1*K
  mvin(IA[...:...]);
  for (n2=0; n2<N2; n2++) {
    // W buffer stores: N1*N0*K
    mvin(W[...:...]);
    OA[...:...]=0;
    for (m1=0; m1<M1; m1++) {
      for (n1=0; n1<N1; n1++) {
        for (k1=0; k1<K1; k1++) {
          ...
        }
      }
    }
    mvout(OA[...:...]);
  }
}
```



Mapping Space

- » Loop ordering
- » Loop bound
- » Spatial choice



Mapping Dimensions

⊙ Loop ordering:

- ◆ Which index goes to the inner/outer loop?



⊙ Loop bounds:

- ◆ What are the loop bounds (i.e., $N?$ / $K?$ / $M?$) for each loop?

⊙ Spatial choice:

- ◆ Which loop should be spatial/temporal?
 - ▣ Data/Model Parallelism

Mapping Problem Example

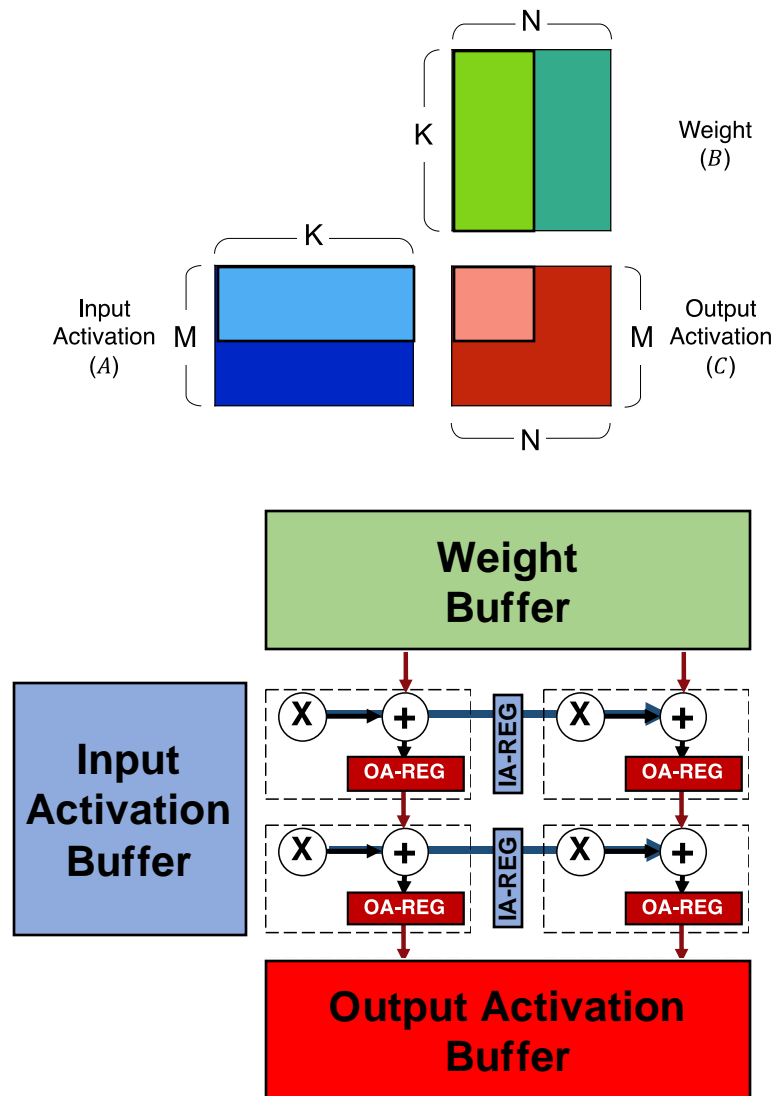


```

for (m2=0; m2<M2; m2++) {
    // IA buffer:  $M1 \times K > 1 \times K$ 
    mvin(IA[m2*M1:(m2+1)*M1,0:K]);
    for (n2=0; n2<N2; n2++) {
        // W buffer:  $N1 \times N0 \times K < K \times N$ 
        mvin(W[0:K,n2*N1*N0:(n2+1)*N1*N0]);
        OA[n2*N1*N0:(n2+1)*N1*N0,m2*M1:(m2+1)*M1]=0;
        for (m1=0; m1<M1; m1++) {
            for (n1=0; n1<N1; n1++) {
                for (k1=0; k1<K1; k1++) {
                    parallel_for (n0=0; n0<N0; n0++) {
                        parallel_for (k0=0; k0<K0; k0++) {
                            OA[n2*N1*N0+n1*N0+n0,m2*M1+m1]
                                += IA[m2*M1+m1,k1*K1+k0]
                                    * W [k1*K0+k0,n2*N1*N0+n1*N0+n0];
                        }
                    }
                }
            }
        }
        mvout(OA[n2*N1*N0:(n2+1)*N1*N0,m2*M1:(m2+1)*M1]);
    }
}

```

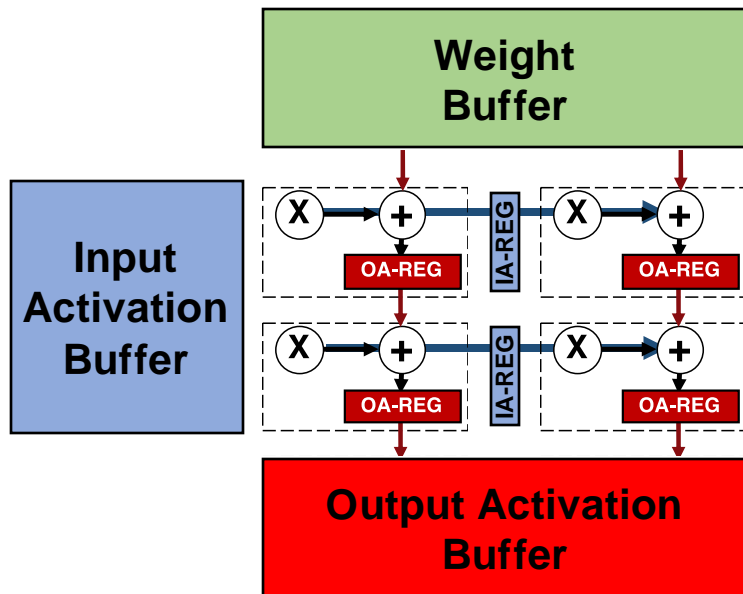

Loop Bound



```

for (m2=0; m2<M2; m2++) {
    // IA buffer: M1*K > 1*K
    mvin(IA[m2*M1:(m2+1)*M1,0:K]);
    for (n2=0; n2<N2; n2++) {
        // W buffer: N1*N0*K < K*N
        mvin(W[0:K,n2*N1*N0:(n2+1)*N1*N0]);
        OA[n2*N1*N0:(n2+1)*N1*N0,m2*M1:(m2+1)*M1]=0;
        for (m1=0; m1<M1; m1++) {
            for (n1=0; n1<N1; n1++) {
                for (k1=0; k1<K1; k1++) {
                    parallel_for (n0=0; n0<N0; n0++) {
                        parallel_for (k0=0; k0<K0; k0++) {
                            OA[n2*N1*N0+n1*N0+n0,m2*M1+m1]
                                += IA[m2*M1+m1,k1*K1+k0]
                                   * W [k1*K0+k0,n2*N1*N0+n1*N0+n0];
                        }
                    }
                }
            }
        }
        compute_matmul();
    }
    mvout(OA[n2*N1*N0:(n2+1)*N1*N0,m2*M1:(m2+1)*M1]);
}
    
```

Loop Ordering (1/2)



```

for (m2=0; m2<M2; m2++) {
    // IA buffer:  $M1 \times K > 1 \times K$ 
    mvin(IA[m2*M1:(m2+1)*M1,0:K]);
    for (n2=0; n2<N2; n2++) {
        // W buffer:  $N1 \times N0 \times K < K \times N$ 
        mvin(W[0:K,n2*N1*N0:(n2+1)*N1*N0]);
        compute_matmul(*W, *IA, *OA,
                        N2, N1, N0,
                        M2, M1,
                        K1, K0,
                        m2, n2);
        mvout(OA[n2*N1*N0:(n2+1)*N1*N0,
                 m2*M1:(m2+1)*M1]);
    }
}

```



Loop Ordering (2/2)

Option 1: Loops $m2 \rightarrow n2$

```
for (m2=0; m2<M2; m2++) {  
    // IA buffer: M1*K  
    mvin(IA[m2*M1:(m2+1)*M1,0:K]);  
    for (n2=0; n2<N2; n2++) {  
        // W buffer: N1*N0*K  
        mvin(W[0:K,n2*N1*N0:(n2+1)*N1*N0]);  
        compute_matmul(*W, *IA, *OA,  
                        ...  
                        m2, n2);  
        mvout(OA[n2*N1*N0:(n2+1)*N1*N0,  
               m2*M1:(m2+1)*M1]);  
    }  
}
```

IA Movement: $M * K$

W Movement: $M2 * N * K$

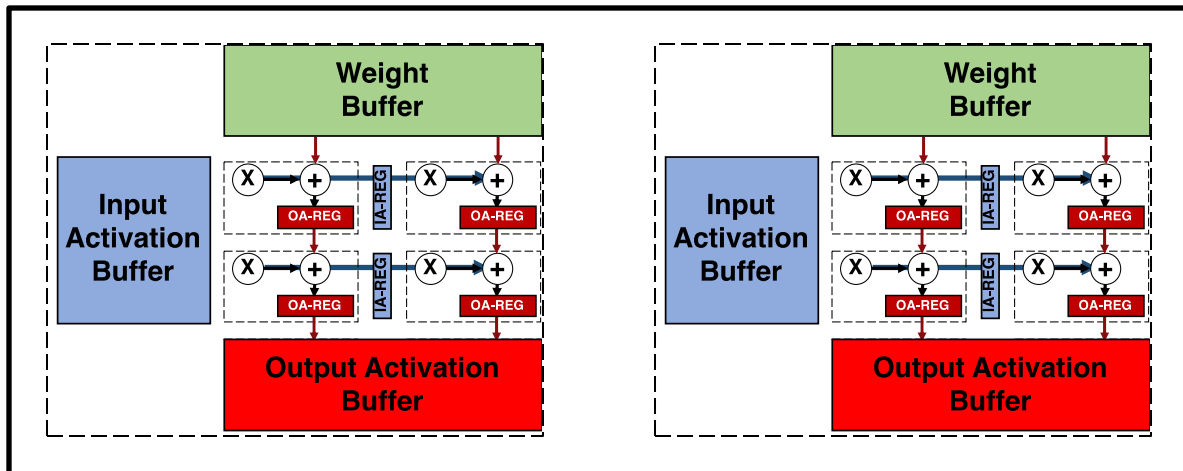

Option 2: Loops $n2 \rightarrow m2$

```
for (n2=0; n2<N2; n2++) {  
    // W buffer: N1*N0*K  
    mvin(W[0:K,n2*N1*N0:(n2+1)*N1*N0]);  
    for (m2=0; m2<M2; m2++) {  
        // IA buffer: M1*K  
        mvin(IA[m2*M1:(m2+1)*M1,0:K]);  
        compute_matmul(*W, *IA, *OA,  
                        ...  
                        m2, n2);  
        mvout(OA[n2*N1*N0:(n2+1)*N1*N0,  
               m2*M1:(m2+1)*M1]);  
    }  
}
```

IA Movement: $N2 * M * K$

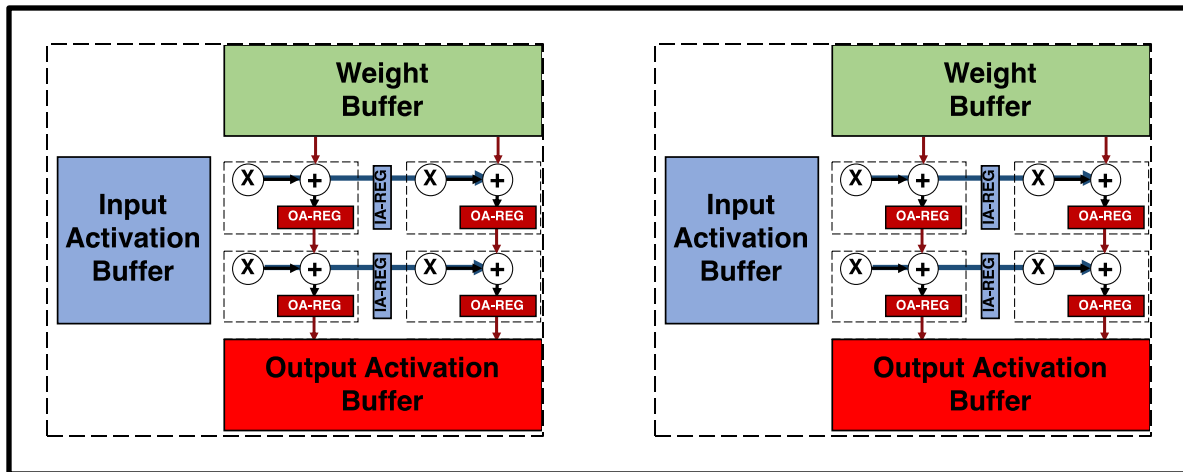

W Movement: $N * K$

Spatial Choice: Model Parallelism



```
parallel_for (n2=0; n2<N2; n2++) {
    // Model Parallelism
    mvin(W[0:K, n2*N1*N0:(n2+1)*N1*N0]);
    for (m2=0; m2<M2; m2++) {
        mvin(IA[m2*M1:(m2+1)*M1, 0:K]);
        compute_matmul(*W, *IA, *OA,
            ...
            m2, n2);
        mvout(OA[n2*N1*N0:(n2+1)*N1*N0,
            m2*M1:(m2+1)*M1]);
    }
}
```

Spatial Choice: Data Parallelism



```
parallel_for (m2=0; m2<M2; m2++) {
    // Data Parallelism
    mvin(IA[m2*M1:(m2+1)*M1,0:K]);
    for (n2=0; n2<N2; n2++) {
        mvin(W[0:K,n2*N1*N0:(n2+1)*N1*N0]);
        compute_matmul(*W, *IA, *OA,
            ...
            m2, n2);
        mvout(OA[n2*N1*N0:(n2+1)*N1*N0,
            m2*M1:(m2+1)*M1]);
    }
}
```



Model Parallelism vs. Data Parallelism

⊙ Data parallelism

- ◆ Each engine processes a subset of the data using the same model
- ◆ Each engine runs the entire model on its assigned data

⊙ Model parallelism

- ◆ The model itself is split across multiple HW engines

⊙ Variants of model parallelism

- ◆ Tensor parallelism
 - Individual tensors are split across multiple engines
 - E.g., splitting weight matrices
 - Aggregation is required during computation
- ◆ Pipeline parallelism
 - The model is divided into stages and executed sequentially in a pipeline
 - Intermediate activations are passed along the pipeline



Tuning of DNN Mapping

Mapping Dimensions

⊙ DNN mapping problem → an **optimization problem**

⊙ Given:

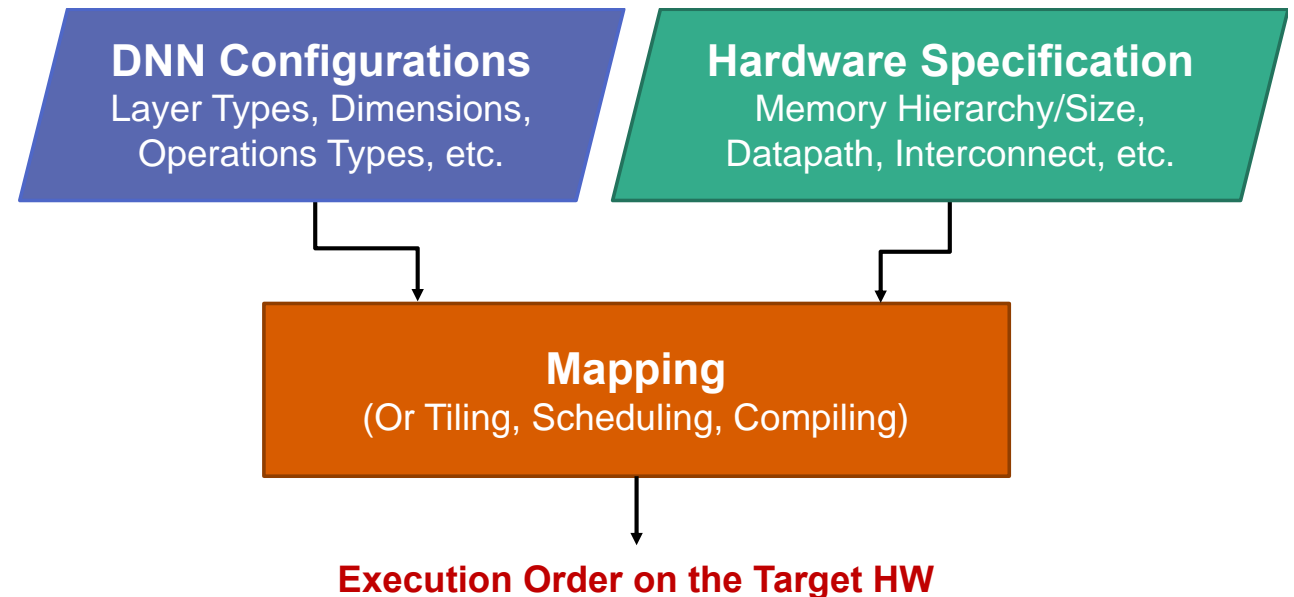
- ◆ DNN dimensions (N, H, W, C, R, S, K, stride, padding)
- ◆ Hardware specifications (dataflow, memory hierarchy)

⊙ Objective:

- ◆ An optimal loop nest that minimizes latency and/or energy
 - ▣ Both temporal and spatial execution order

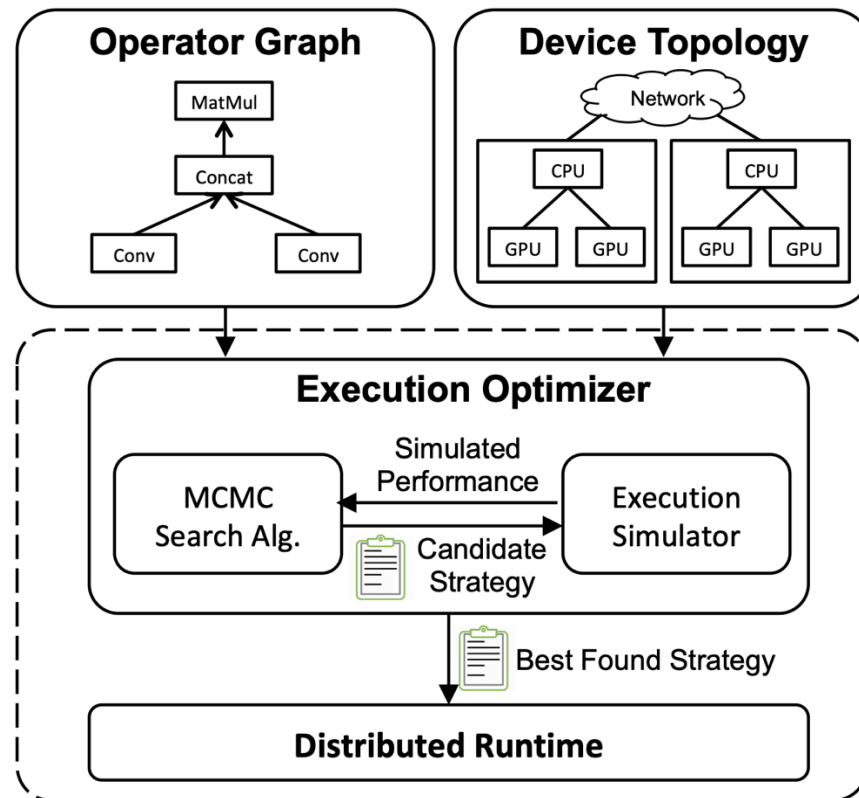
⊙ Approaches:

- ◆ Exhaustive search
- ◆ Random search
- ◆ Learning-based algorithms



Example: FlexFlow, SysML'2018

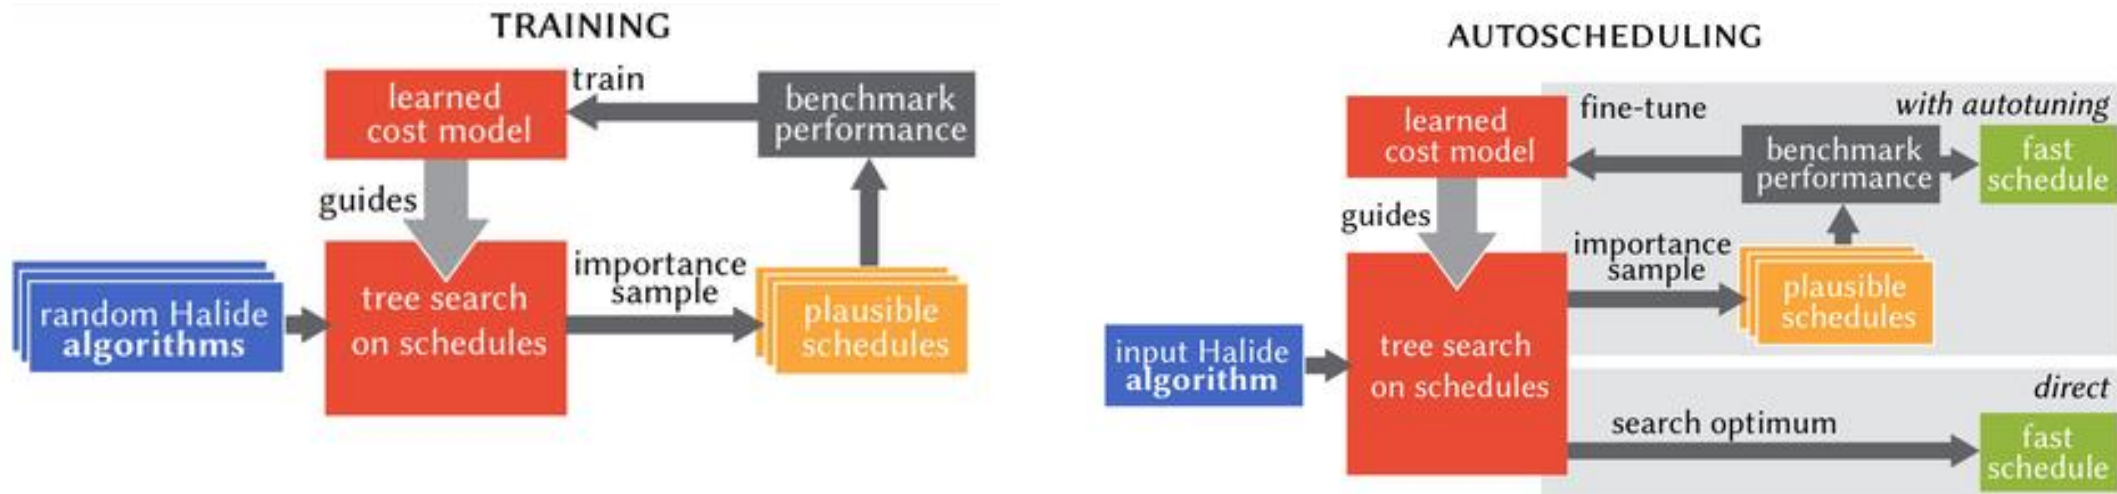
- “The optimizer uses a MCMC (Markov Chain Monte Carlo) search algorithm to explore the space of possible parallelization strategies and iteratively proposes candidate strategies that are evaluated by an execution simulator.”



Beyond Data and Model Parallelism for Deep Neural Networks, SysML 2018

Example: Halide, SIGGRAPH'2019

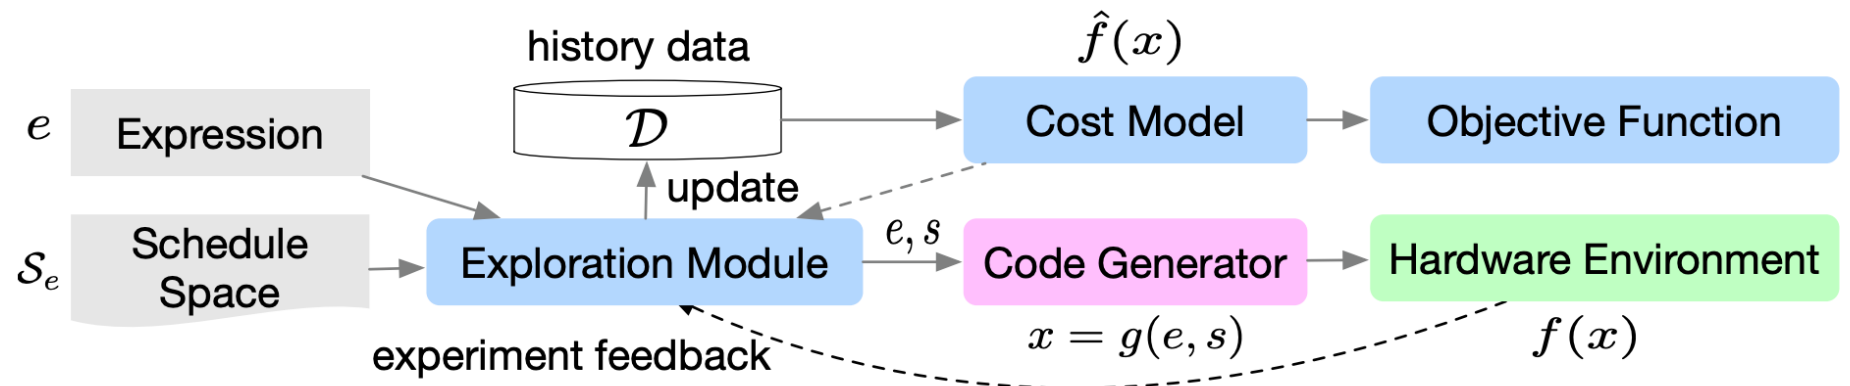
- “We generate schedules for Halide programs using tree search over the space of schedules guided by a learned cost model and optional autotuning. The cost model is trained by benchmarking thousands of randomly-generated Halide programs and schedules. The resulting code significantly outperforms prior work and human experts.”



Learning to Optimize Halide with Tree Search and Random Programs, *SIGGRAPH 2019*

Example: TVM, NeurIPS'2018

- Learn domain-specific statistical cost models to guide the search of tensor operator implementations over billions of possible program variants
- Further accelerate the search using effective model transfer across workloads



Learning to Optimize Tensor Programs, *NeurIPS'2018*



Summary

- ◎ Temporal and Spatial Mapping based on hardware constraints
 - ◆ Memory hierarchy
 - ◆ Parallelism
- ◎ Tile the loops to improve reuse and parallelism
 - ◆ Loop ordering
 - ◆ Loop bounds
 - ◆ Spatial choices
- ◎ Navigate the large mapping space
 - ◆ Finding an optimized solution
 - ◆ Getting to the solution fast enough